

Containerised ASP.NET Core apps with Docker and Kubernetes

Mete Atamel

Developer Advocate @ Google



Mete Atamel

Developer Advocate for Google Cloud

@meteatamel

atamel@google.com

meteatamel.wordpress.com

Please send talk feedback: bit.ly/atamel

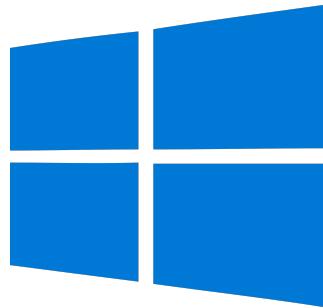


The .NET Revolution

The world of 2014

@meteatamel

Visual Studio



ASP.NET

C#

SQL Server

PowerShell

Eclipse

MySQL

Java

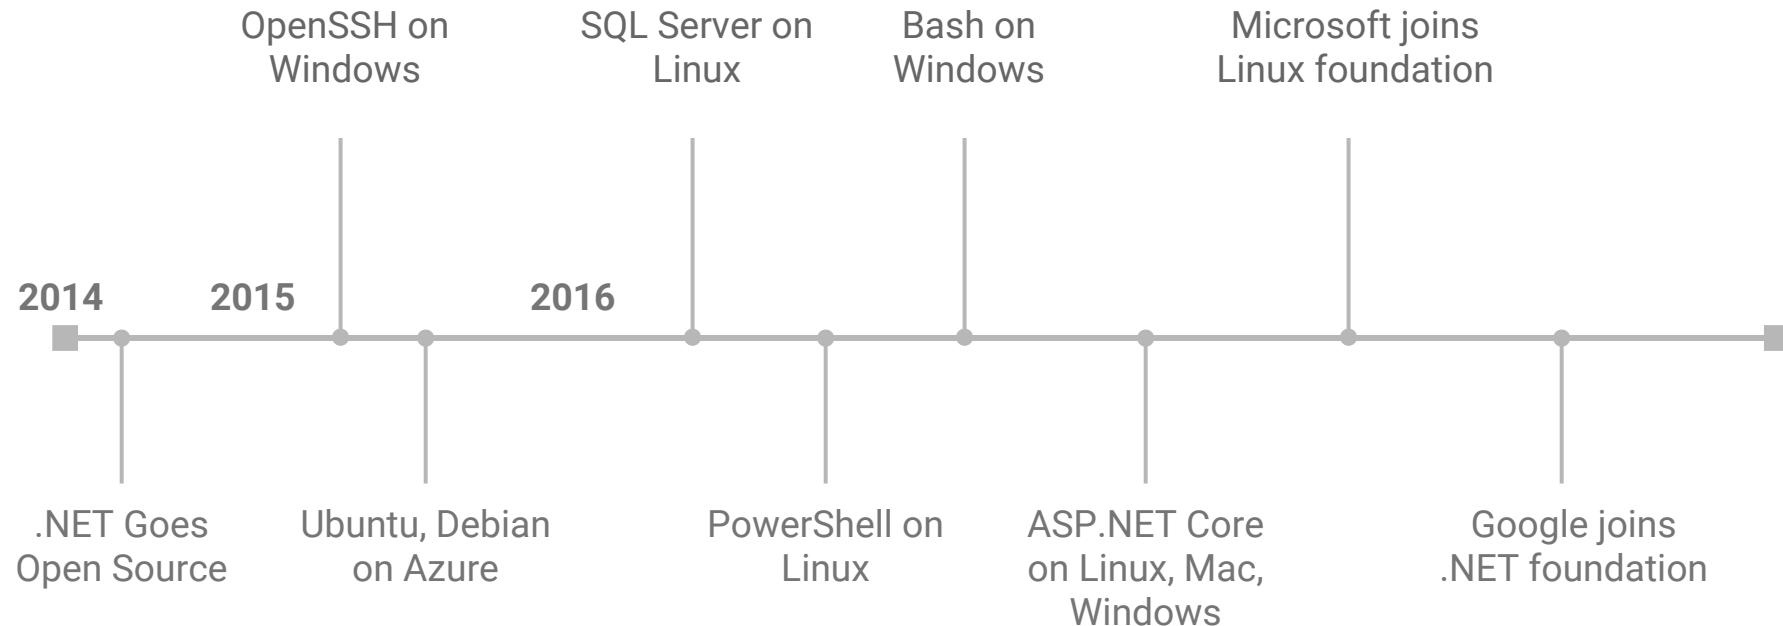
Apache



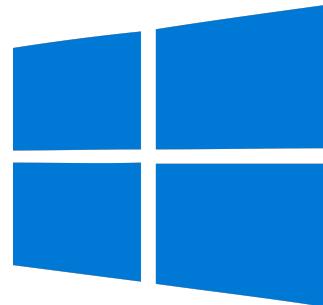
Bash

Things are changing

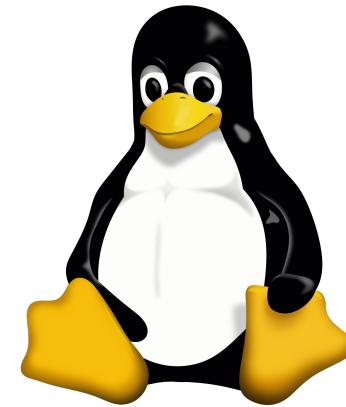
@meteatamel



The convergence



SQL Server Eclipse
Visual Studio Bash
C# Java
ASP.NET PowerShell
MySQL Apache



Great time to be a .NET developer!



However, software development is HARD!

And it is not getting any easier...

In the good old days

@meteatamel

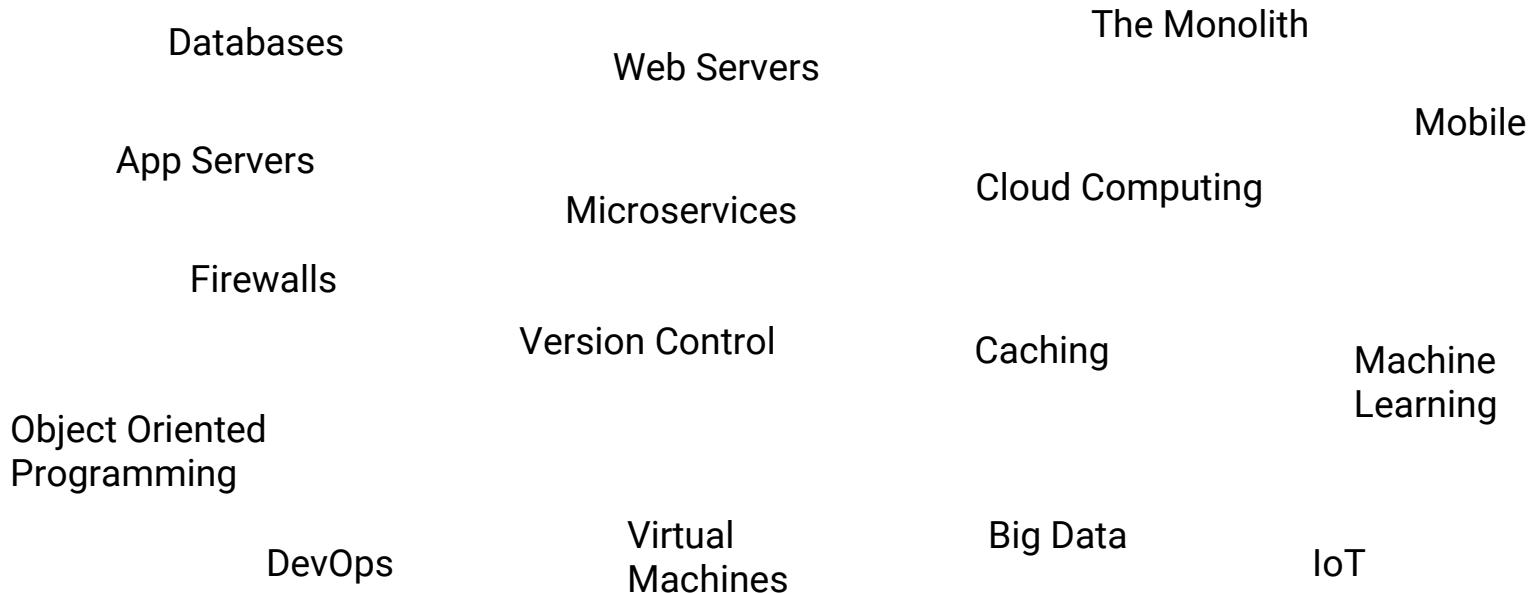
This is all I had to care...

Life was good :-)

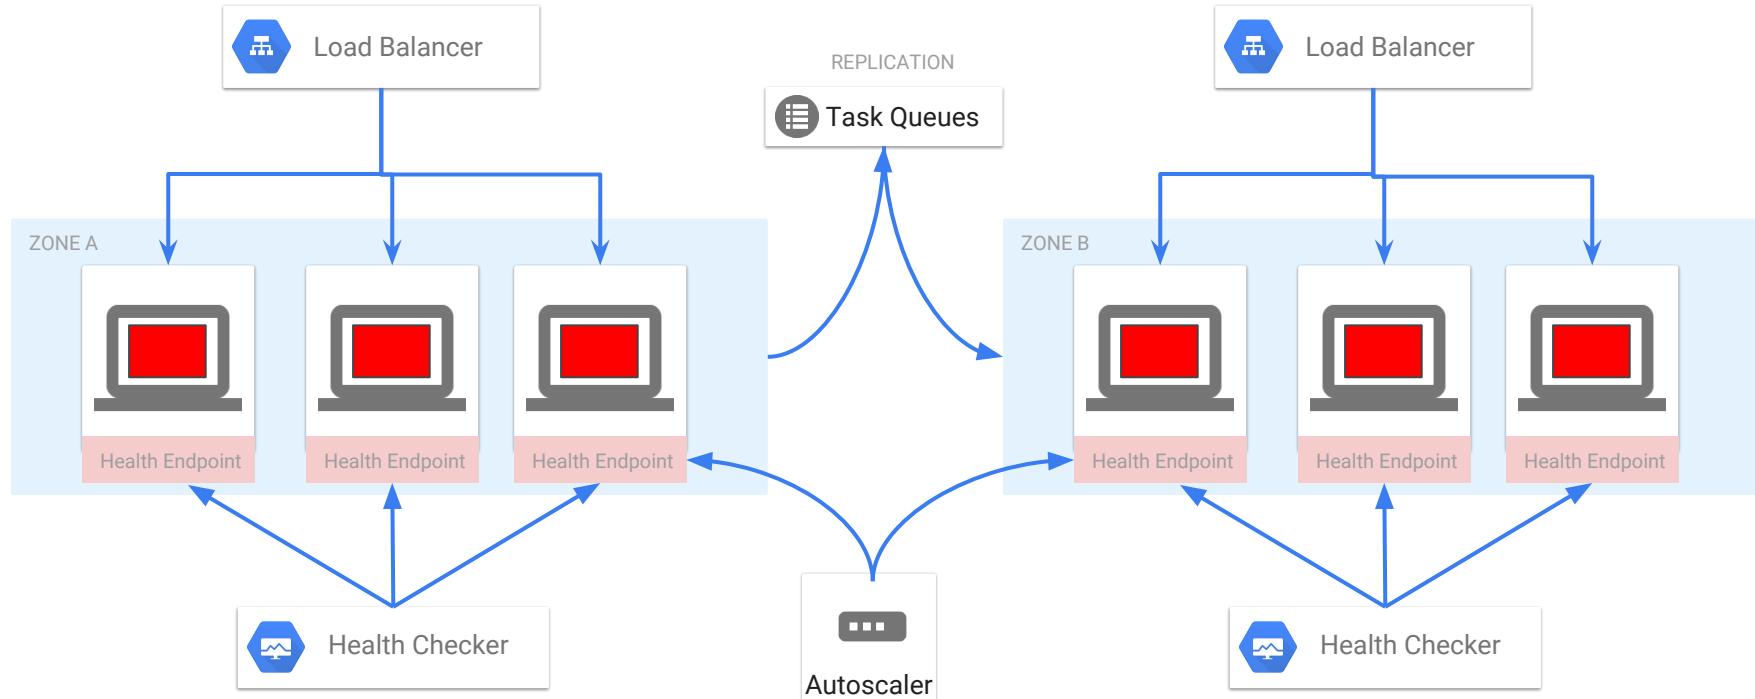


A lot happened since then

Internet



Nowadays



We haven't even talked about

Maintaining code in **different** languages on **different** types of machines

Rolling out the **new version** of your code reliably

Rolling back to the **old version** if something goes wrong

Managing **configuration** and **secrets**

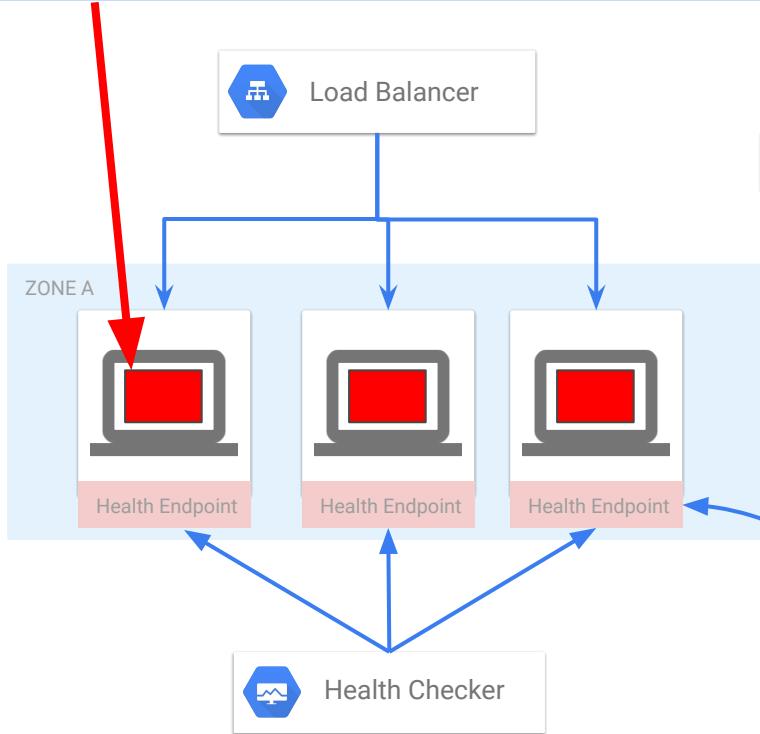
Managing **scripts** that need to run on each machine



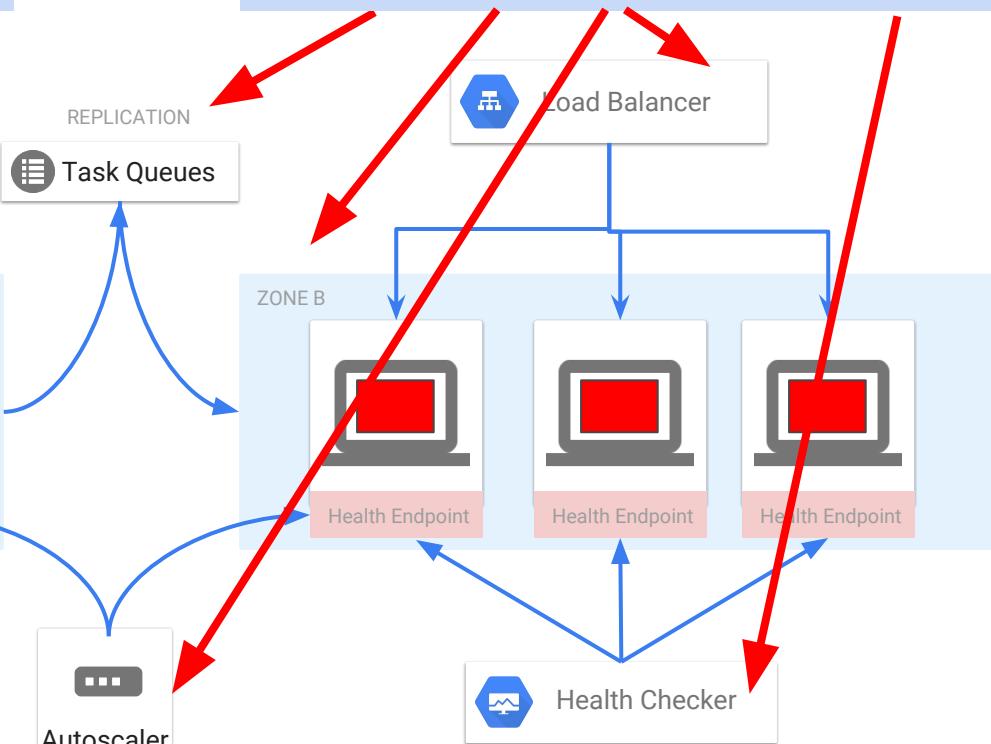
I just want to write some code to solve a real-world problem

The reality

This is all I **want** to care



This is all I **have** to care



Writing code to solve problems is still fun!

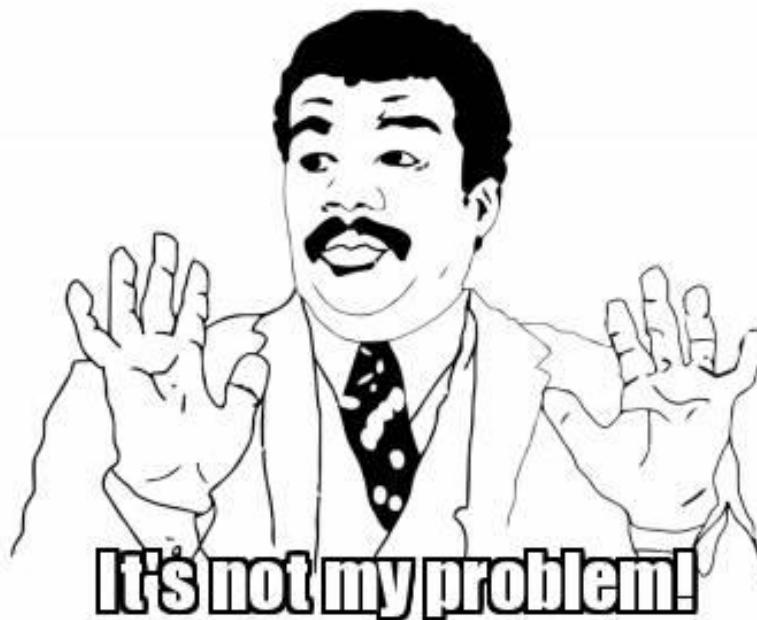
Running that code in production is very hard

What do we do?

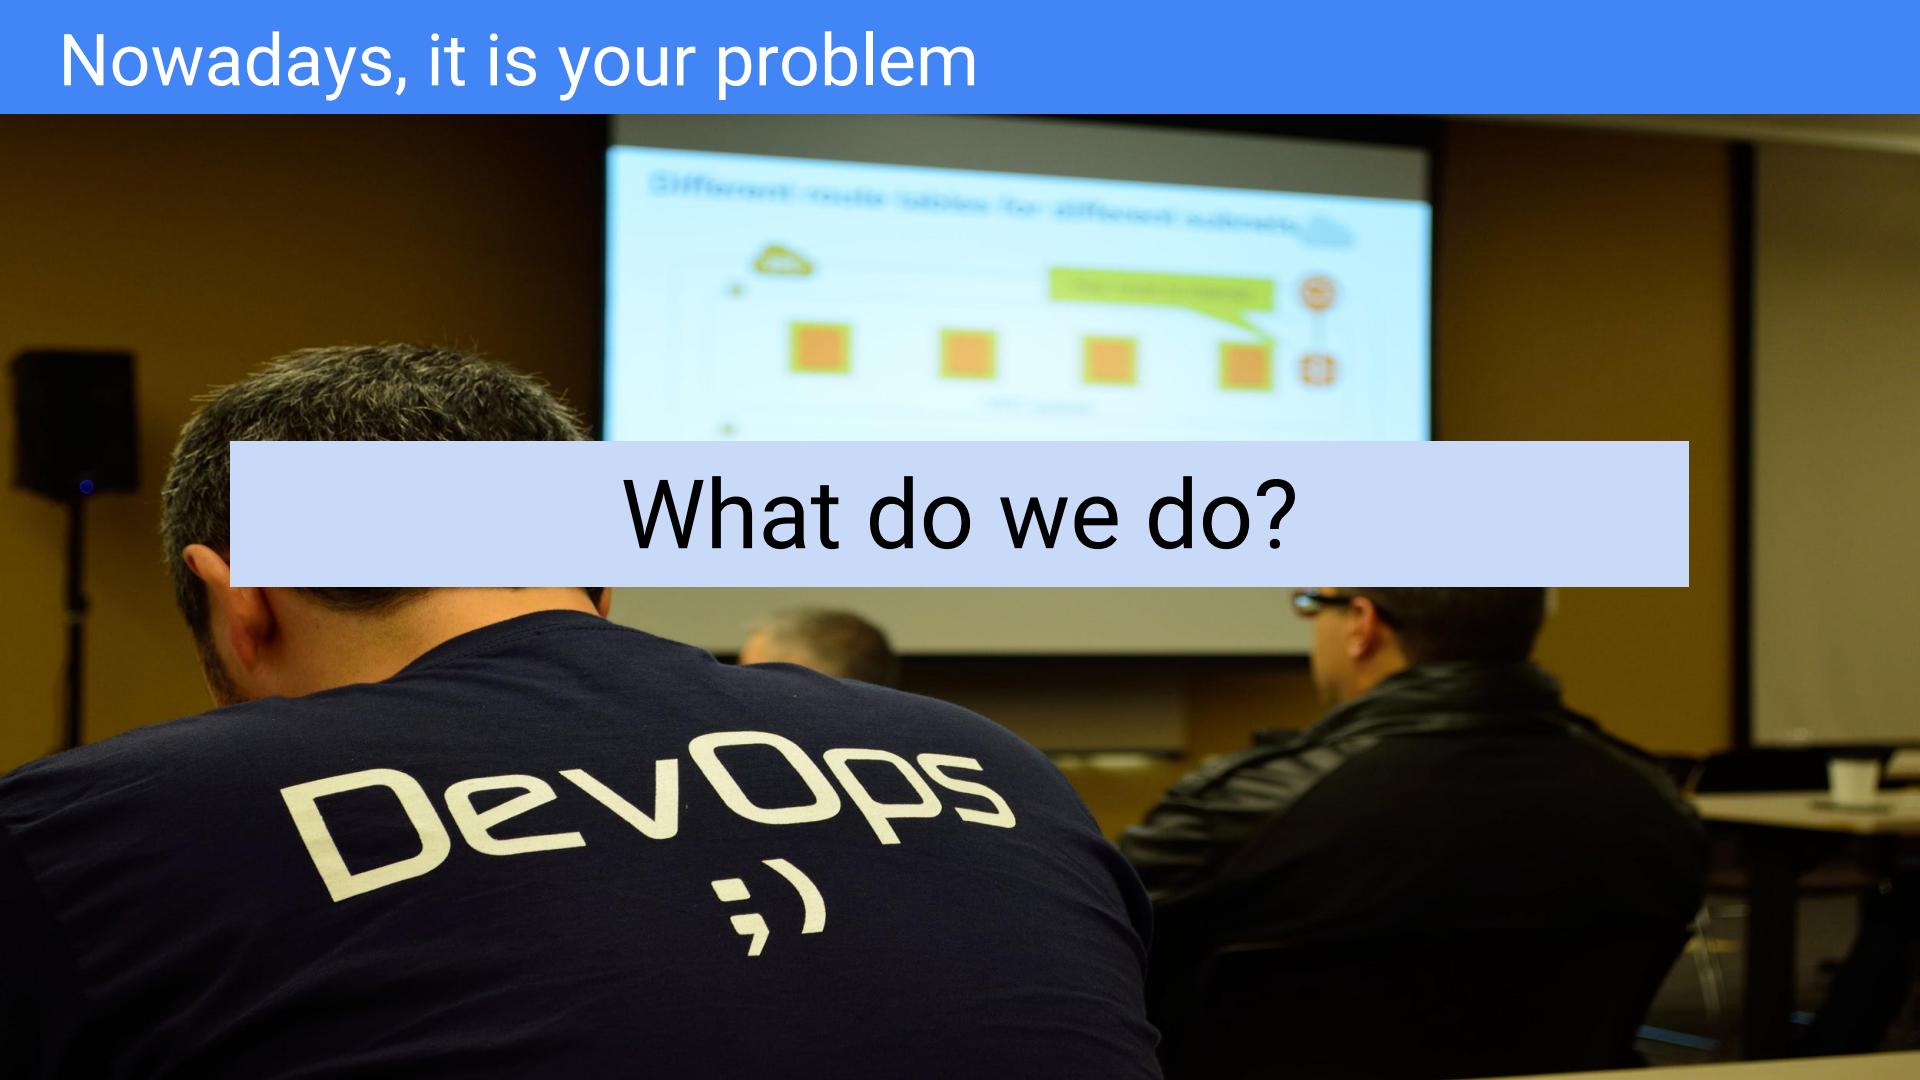
In the good old days

@meteatamel

Write your code, pass it to QA for testing, let operations team run it...



Nowadays, it is your problem

A photograph of a person from behind, wearing a dark blue t-shirt with "DevOps" and a smiley face emoji printed on the back. They are standing in front of a projection screen displaying a slide with various icons. A white rectangular box with a black border is overlaid on the slide, containing the text "What do we do?".

What do we do?

You can write your code in **any** language and run
anywhere exactly the **same** way

Your app is optimally deployed **somewhere** and
managed by **someone**. It just works!

There are no machines. All resources are
automatically provisioned on demand

Docker + Kubernetes + Cloud

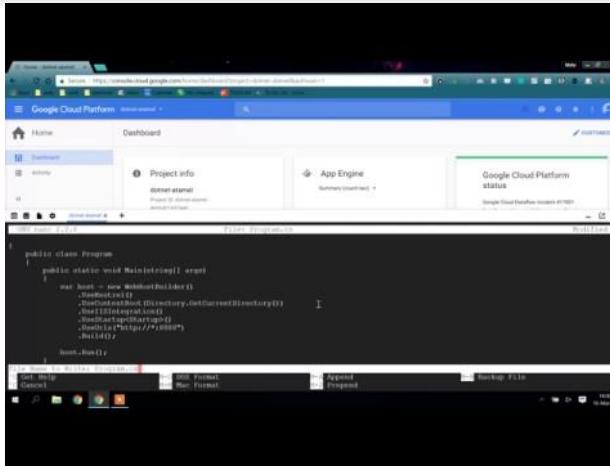
Write your code in any language and run it anywhere exactly the same way
⇒ Containers (eg. Docker, Rkt)

Your app is optimally deployed and managed
⇒ Container Management Platforms (eg. Kubernetes, Docker Swarm, Mesos)

All the resources needed for your app is automatically provisioned per demand
⇒ Cloud Providers (eg. Google Cloud, AWS, Azure)



Demo: Simple Microservice



Containers



What is a container?

@meteatamel

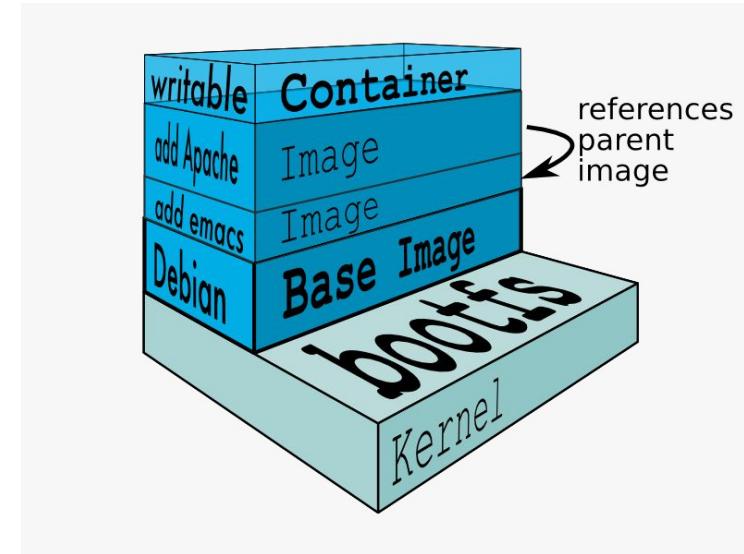
A **lightweight** way to virtualize applications

Linux (or Windows) processes

Lightweight
Hermetically sealed
Isolated

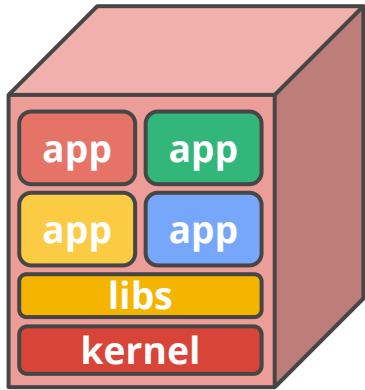
Easily deployable
Introspectable
Composable

Docker



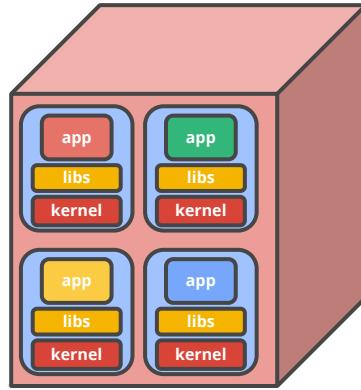
Why containers?

@meteatamel



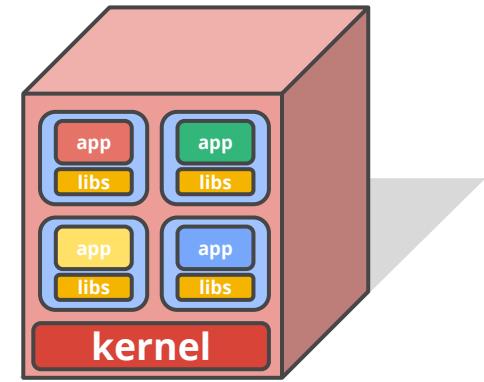
Physical Machine

- ✗ No isolation
- ✗ Common libs
- ✗ Highly coupled Apps & OS



Virtual Machines

- ✓ Isolation
- ✓ No Common Libs
- ✗ Expensive and Inefficient
- ✗ Hard to manage



Containers

- ✓ Isolation
- ✓ No Common Libs
- ✓ Less overhead
- ✗ Less Dependency on Host OS



Google has been developing
and using containers to
manage our applications for
over 12 years.

Everything at Google runs in containers

Gmail, Web Search, Maps, ...

MapReduce, batch, ...

GFS, Colossus, ...

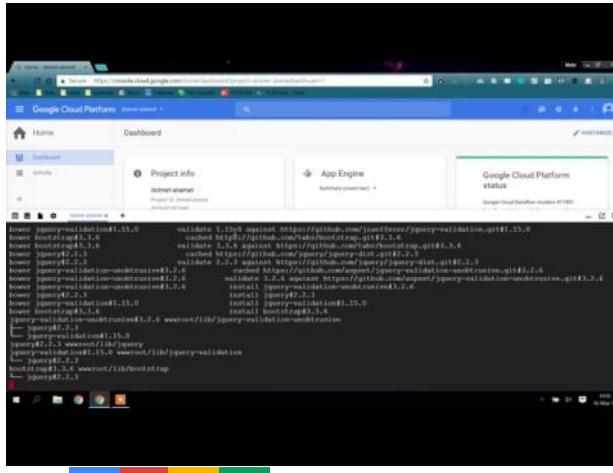
Even **Google's Cloud Platform**: our VMs run in containers!

We launch over **2 billion** containers **per week**



Shipping Containers At Clyde, by Steve Gibson

Demo: Containerised Microservice



Containers help to create a lightweight and consistent environment for apps

But you still need to answer these questions:

- Who takes care of **redundancy**?
- Who takes care of **resiliency**?
- Who **scales up/down** your app?
- Who and how a **new version** of your app gets deployed?
- Who rolls back to a **previous version** if something goes wrong?
- Etc. etc. etc.



Kubernetes



Greek for “*Helmsman*”; also the root of the words “governor” and “cybernetic”

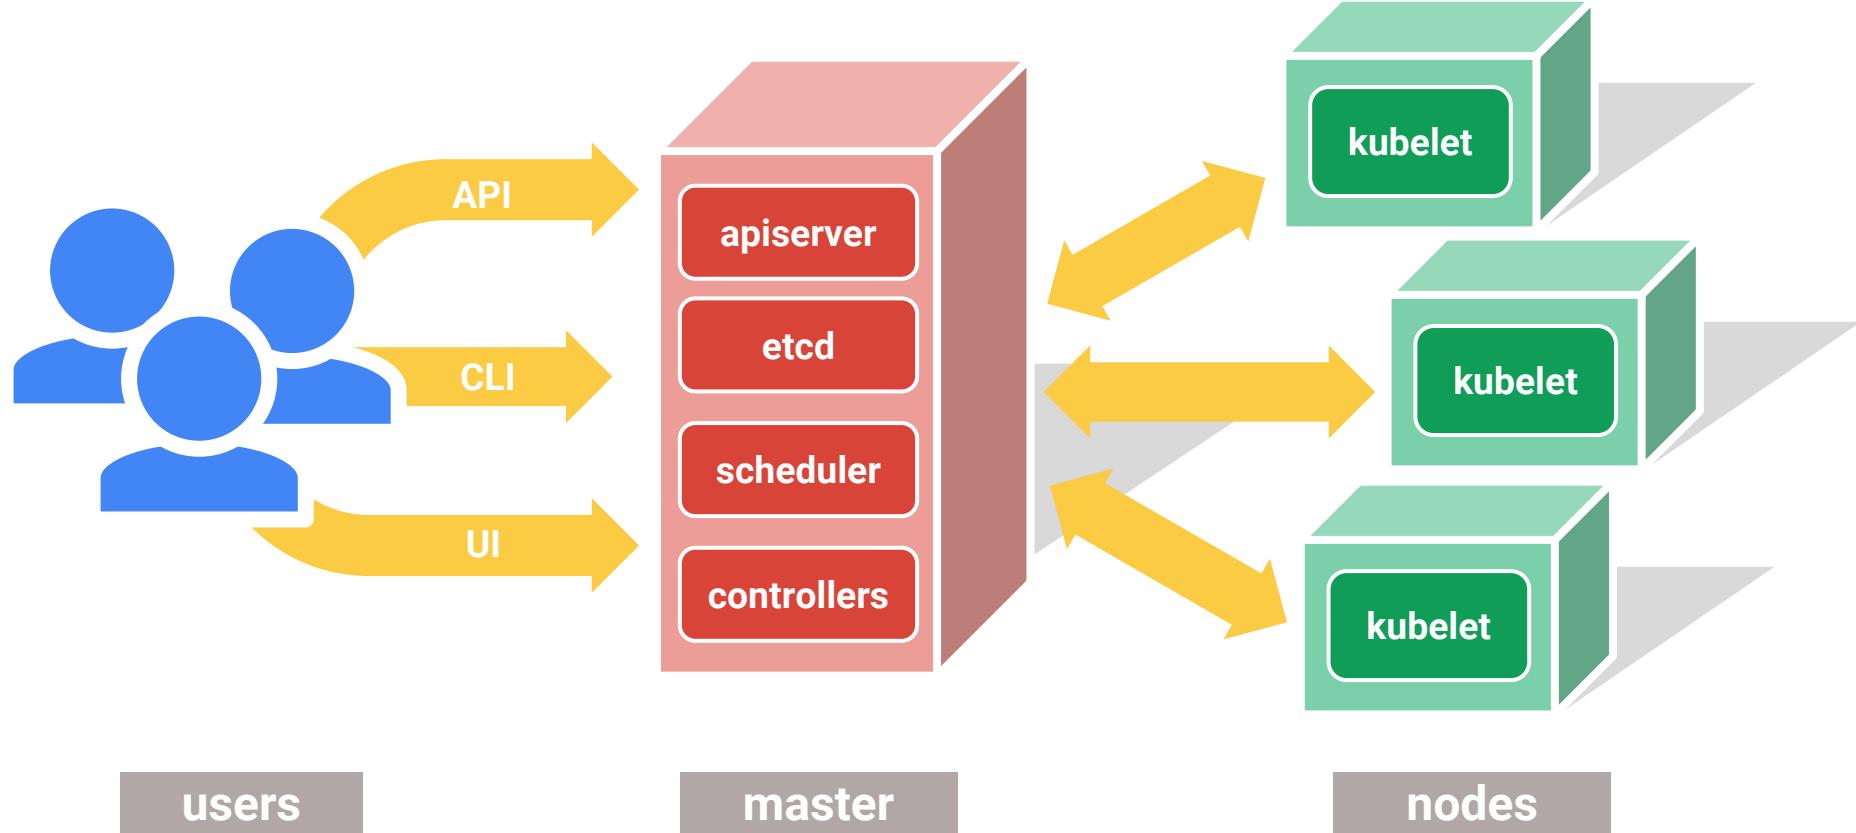
- Manages container clusters
- Inspired and informed by Google’s experiences and internal systems (borg)
- Supports multiple cloud and bare-metal environments
- Supports multiple container runtimes
- **100% Open source**, written in Go

Manage applications, not machines



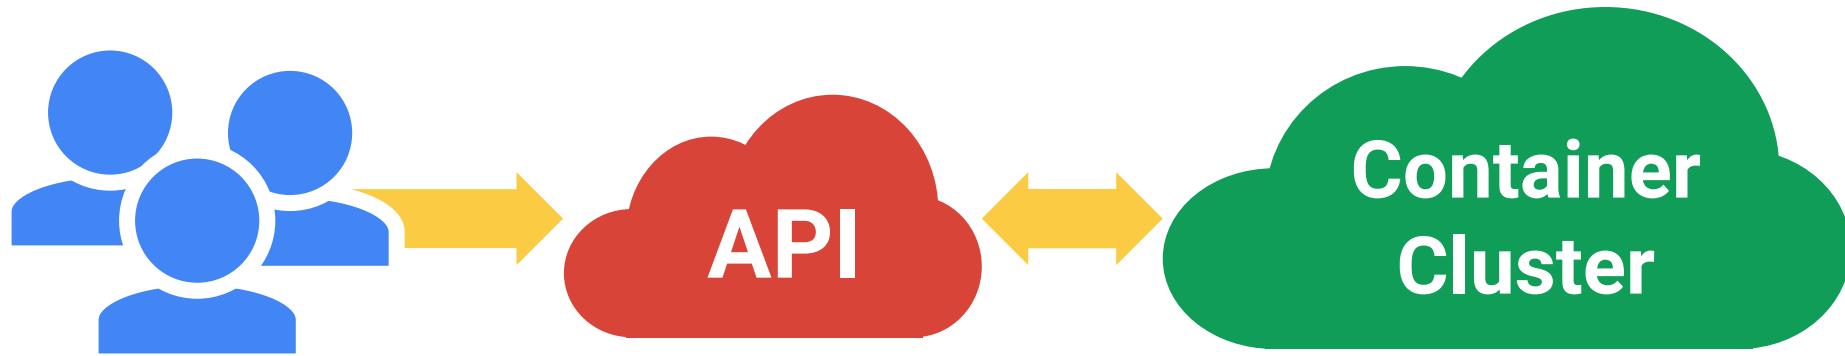
The 10000 foot view

@meteatamel



All you really care about

@meteatamel



1. Setting up the cluster

- Choose a cloud: GCE, AWS, Azure, Rackspace, on-premises, ...
- Choose a node OS: CoreOS, Atomic, RHEL, Debian, CentOS, Ubuntu, ...
- Provision machines: Boot VMs, install and run kube components, ...
- Configure networking: IP ranges for Pods, Services, SDN, ...
- Start cluster services: DNS, logging, monitoring, ...
- Manage nodes: kernel upgrades, OS updates, hardware failures...

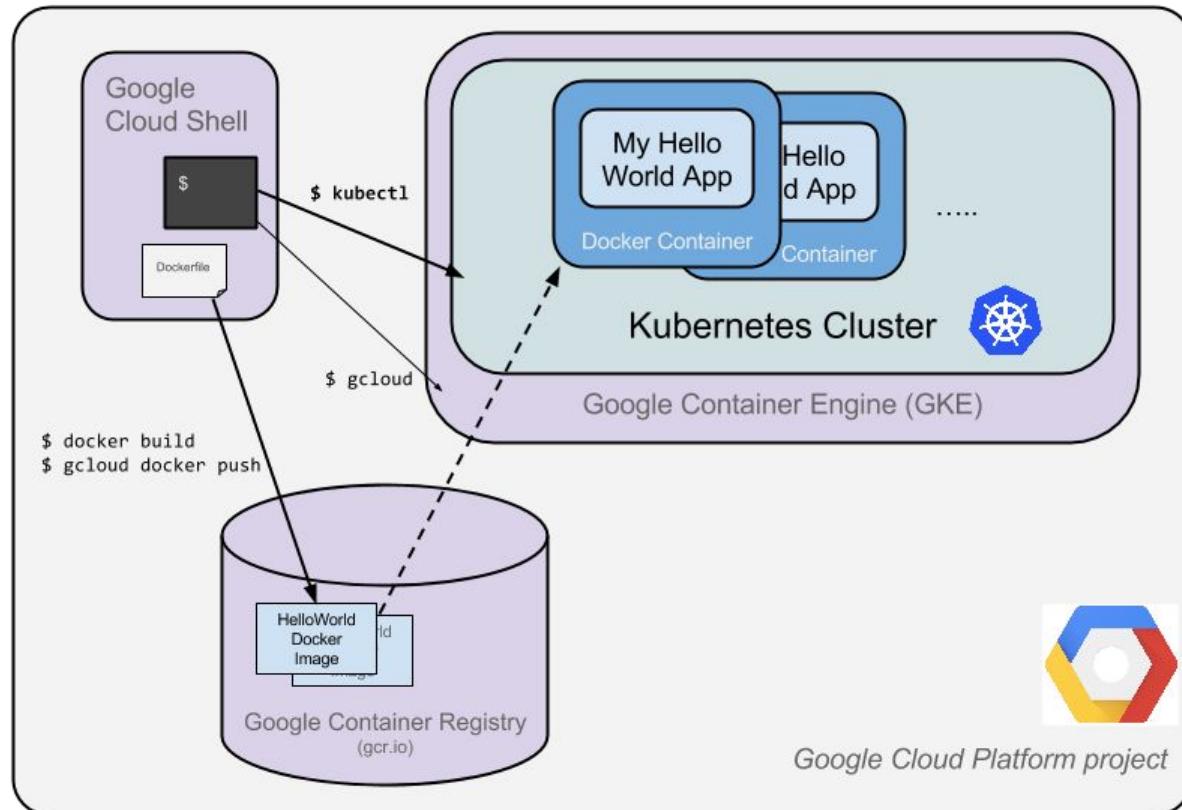
Not the easy or fun part, but unavoidable

This is where things like **Google Container Engine (GKE)** really help

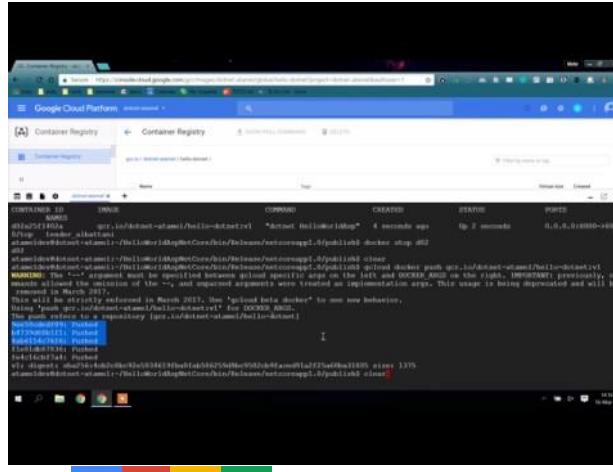


Kubernetes cluster on GKE

@meteatamel



Demo: Create Kubernetes cluster



2. Using the cluster

- Run Pods & Containers
- Replica Sets
- Services
- Volumes

This is the fun part!

A distinct set of problems from cluster setup and management

Don't make developers deal with cluster administration!

Accelerate development by focusing on the applications, not the cluster



Kubernetes Building Blocks



Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

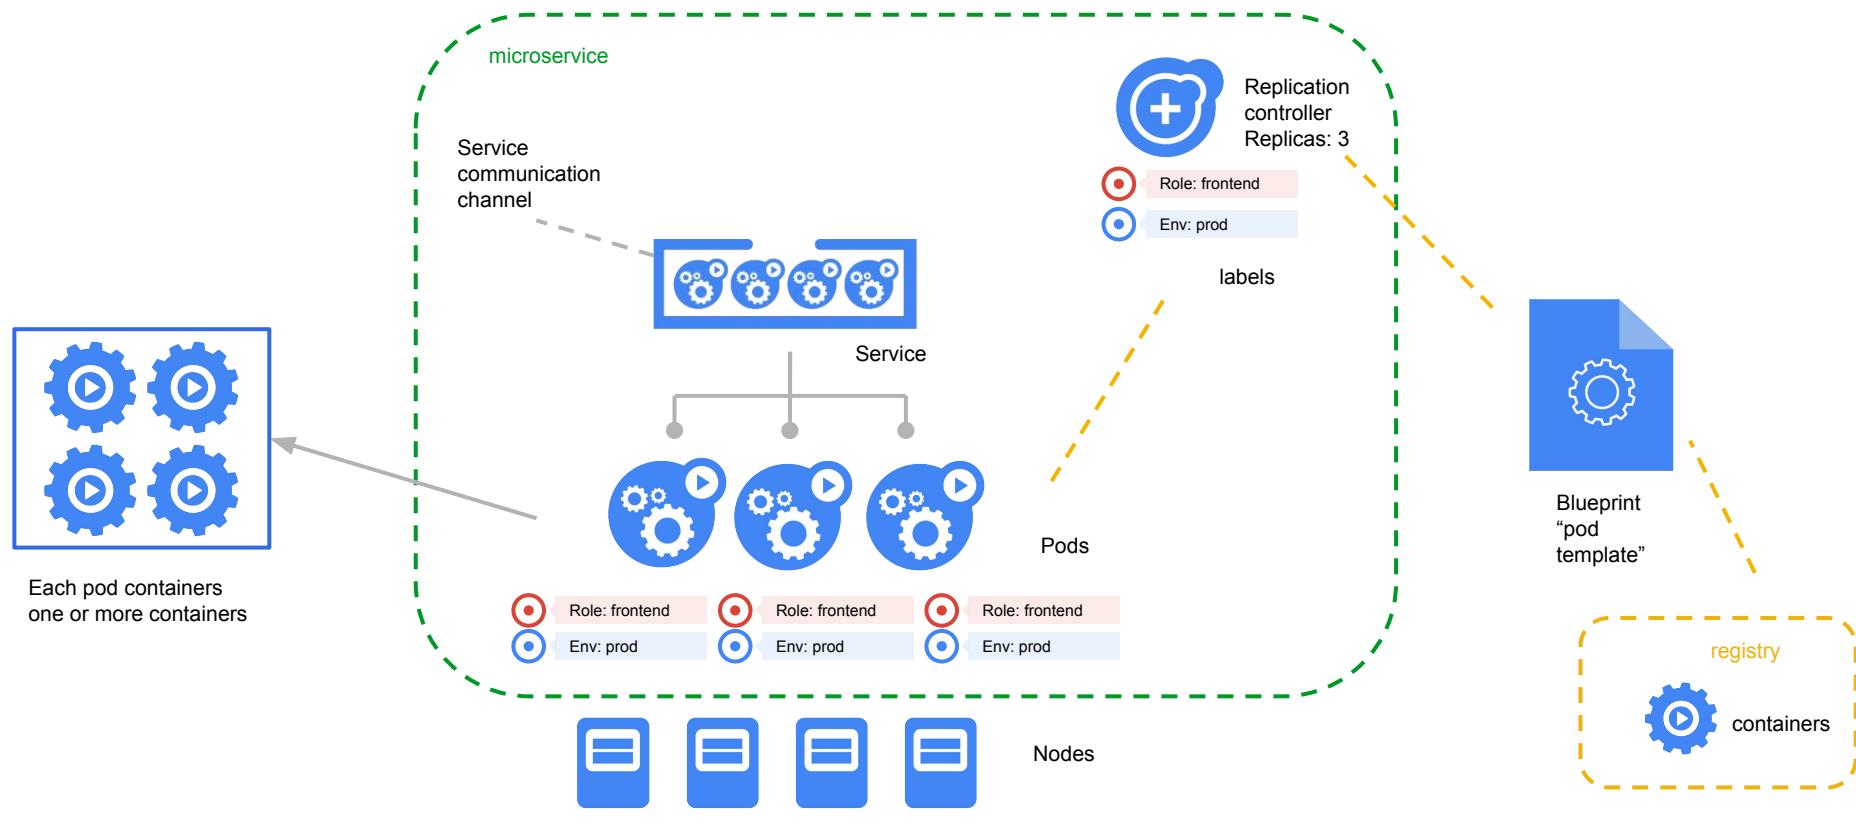
ConfigMap

Selector

Secret



Container cluster



Deployments



Deployments

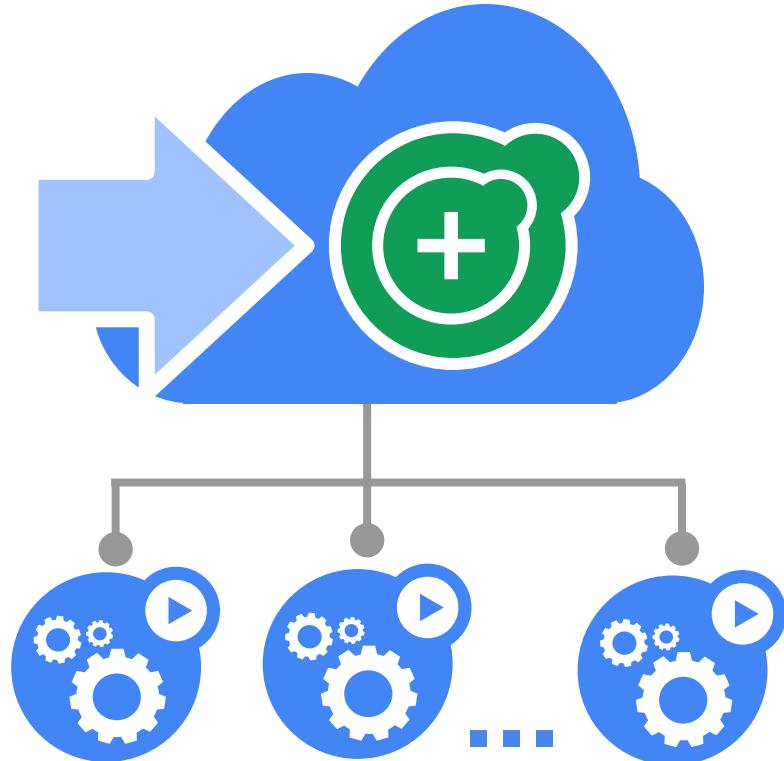
@meteatamel

A Deployment provides declarative updates for Pods and Replica Sets

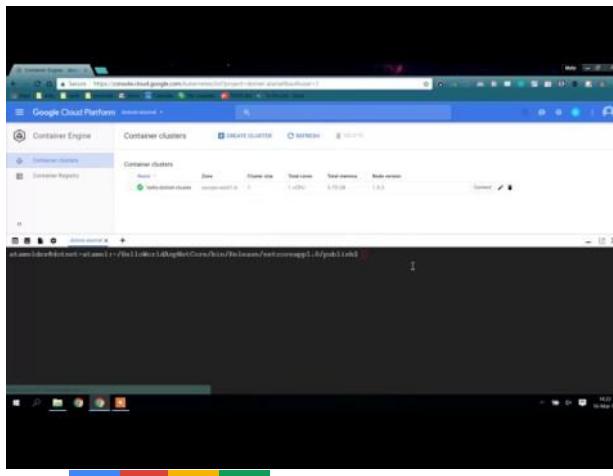
Describe the desired state and the Deployment controller will change the actual state to the desired state at a controlled rate for you.

Deployment manages replica changes for you

- stable object name
- updates are configurable, done server-side
- kubectl edit or kubectl apply



Demo: Create Deployment



Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

ConfigMap

Selector

Secret



Pods and Volumes



Small group of containers & volumes

Tightly coupled

The atom of scheduling & placement

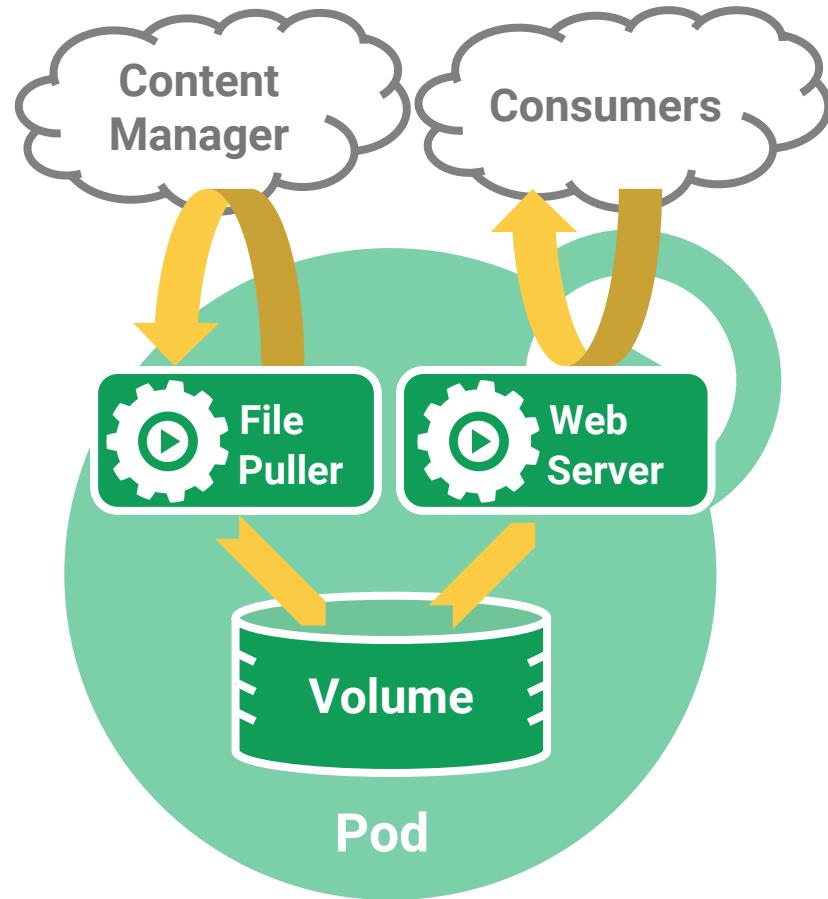
Shared namespace

- share IP address & localhost
- share IPC, etc.

Managed lifecycle

- bound to a node, restart in place
- can die, cannot be reborn with same ID

Example: data puller & web server



Pod-scoped storage

Support many types of volume plugins

- Empty dir (and tmpfs)
- Host path
- Git repository
- GCE Persistent Disk
- AWS Elastic Block Store
- Azure File Storage
- iSCSI
- Flocker
- NFS
- vSphere
- GlusterFS
- Ceph File and RBD
- Cinder
- FibreChannel
- Secret, ConfigMap, DownwardAPI
- Flex (exec a binary)
- ...



Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

ConfigMap

Selector

Secret



Labels & Selectors



Arbitrary metadata

Attached to any API object

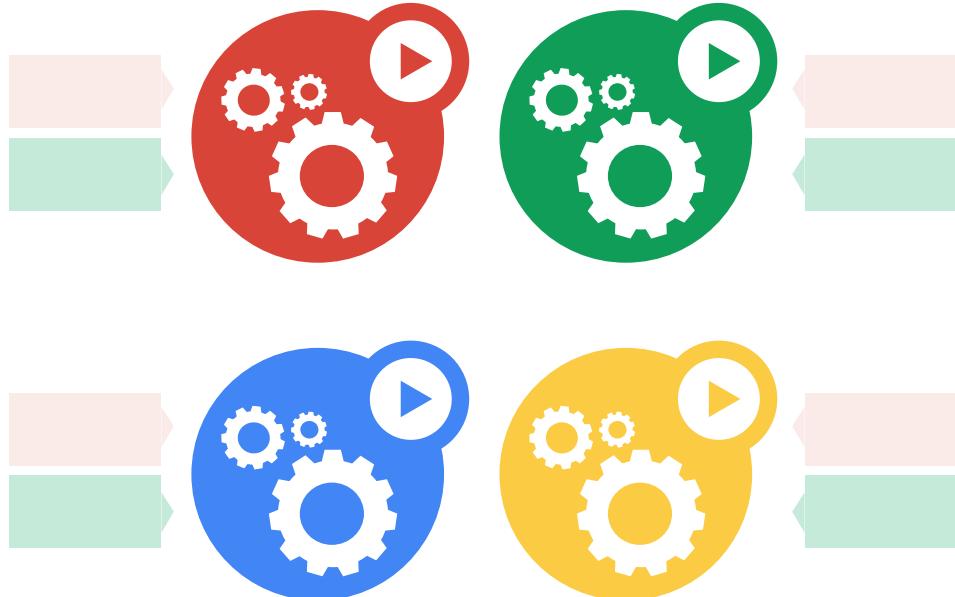
Generally represent **identity**

Queryable by **selectors**

- think SQL 'select ... where ...'

The **only** grouping mechanism

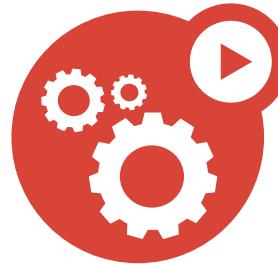
- pods under a ReplicationController
- pods in a Service
- capabilities of a node (constraints)



Selectors

@meteatamel

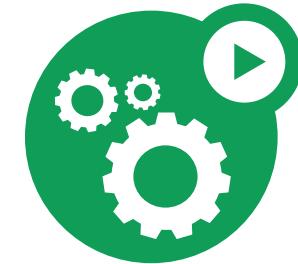
App: MyApp
Phase: prod
Role: FE



App: MyApp
Phase: test
Role: FE



App: MyApp
Phase: prod
Role: BE

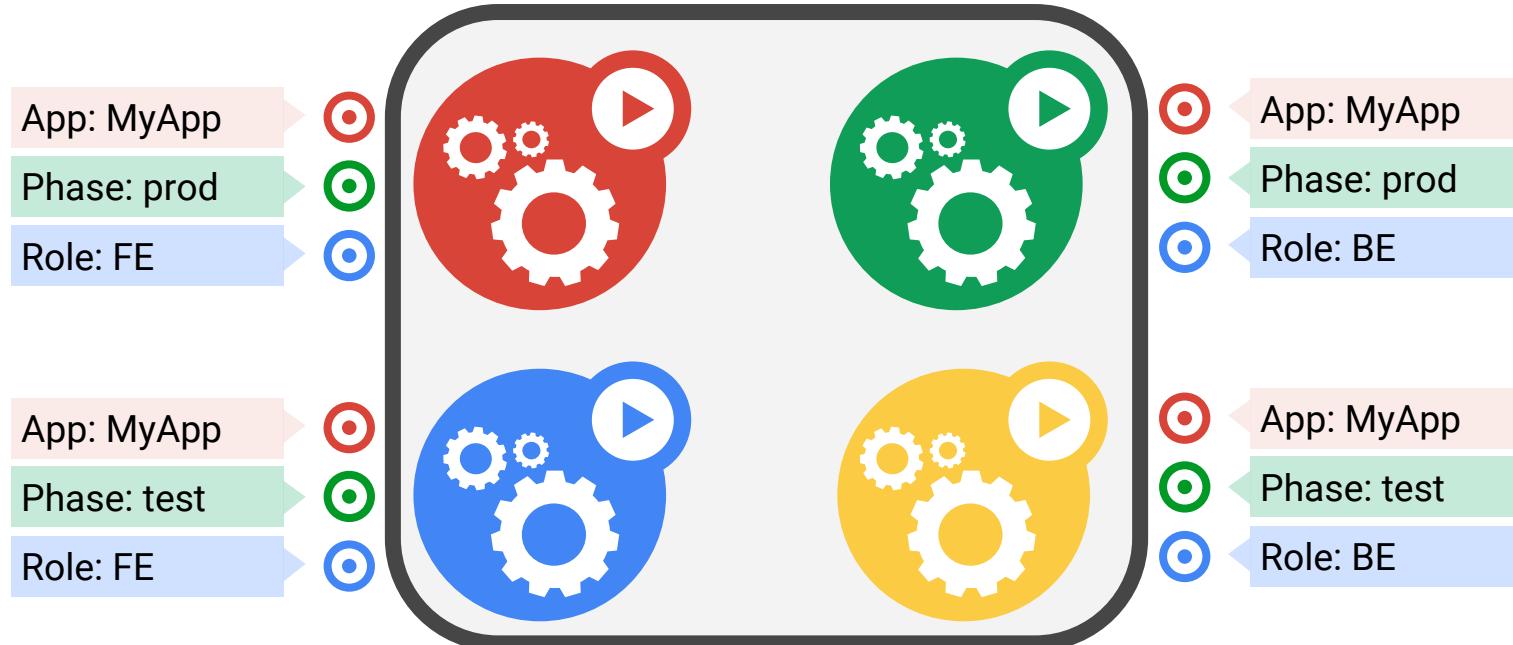


App: MyApp
Phase: test
Role: BE



Selectors

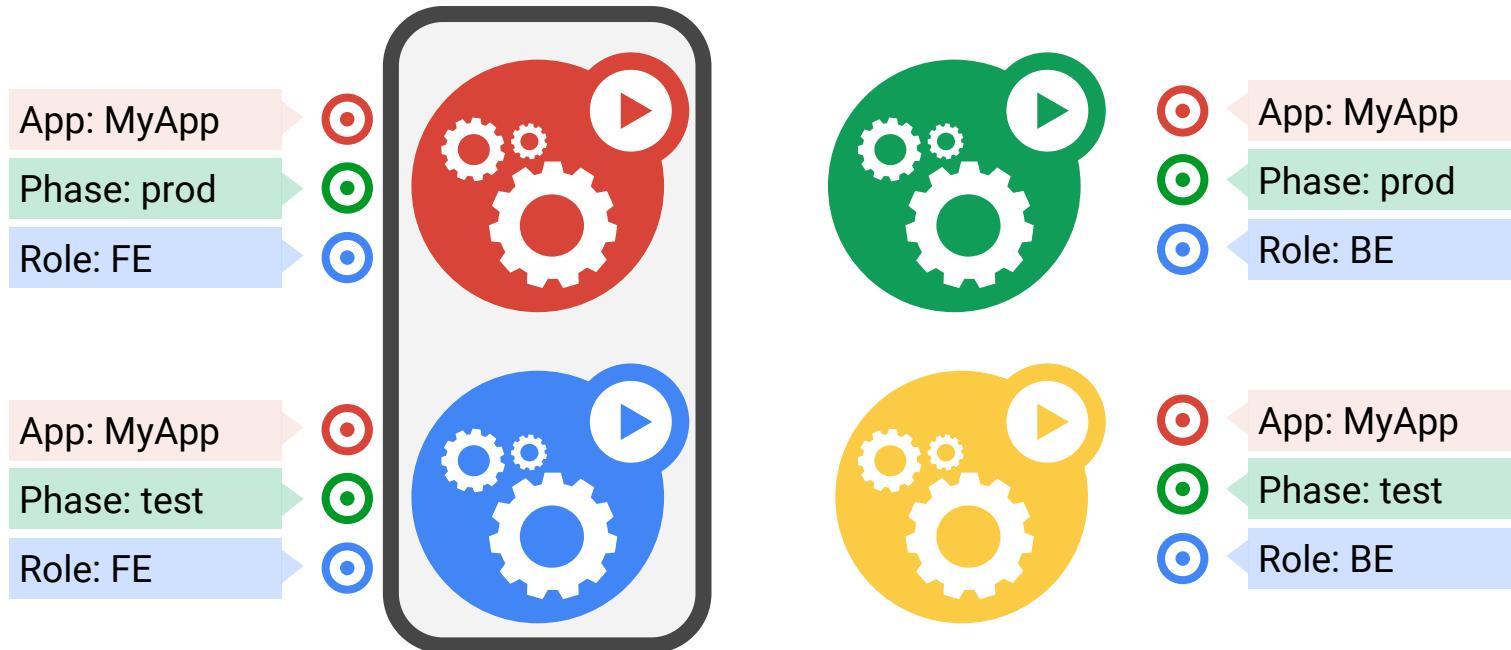
@meteatamel



App = MyApp

Selectors

@meteatamel

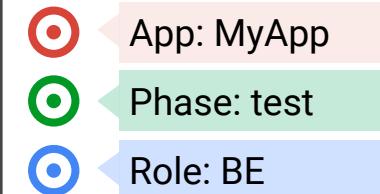
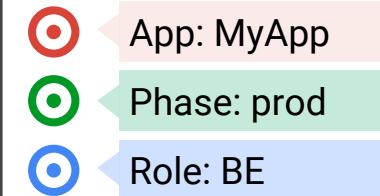
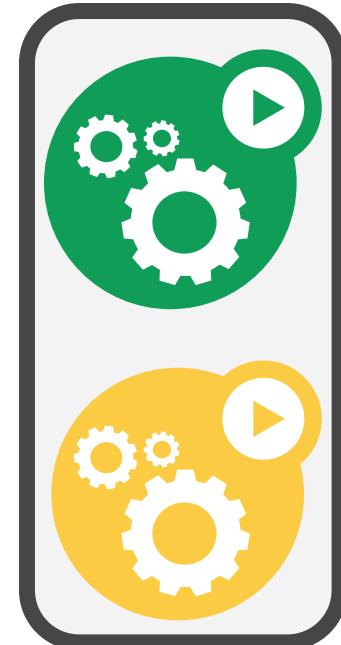
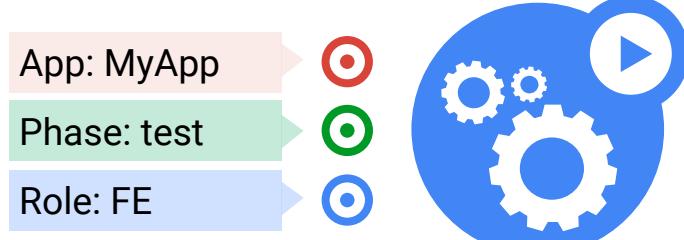
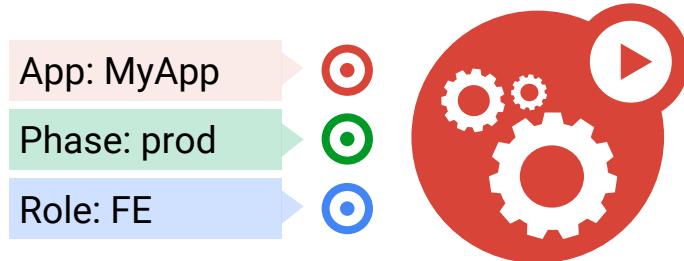


App = MyApp, Role = FE



Selectors

@meteatamel

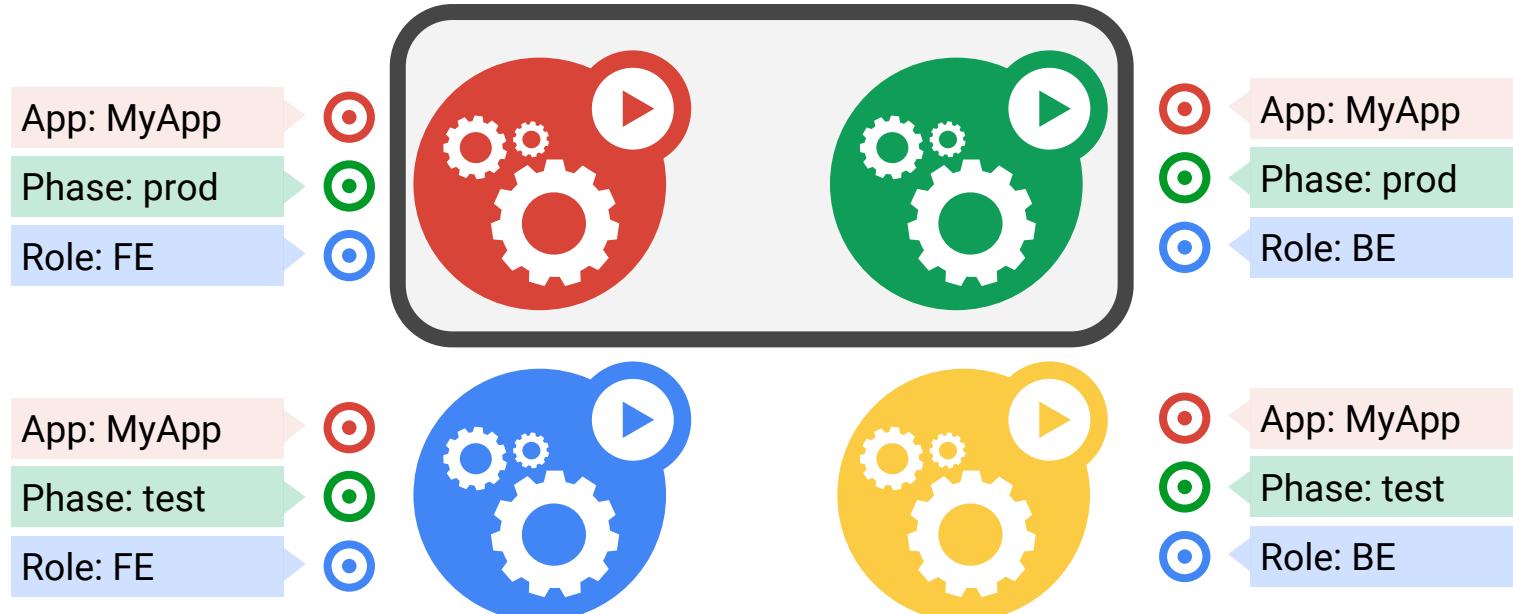


App = MyApp, Role = BE



Selectors

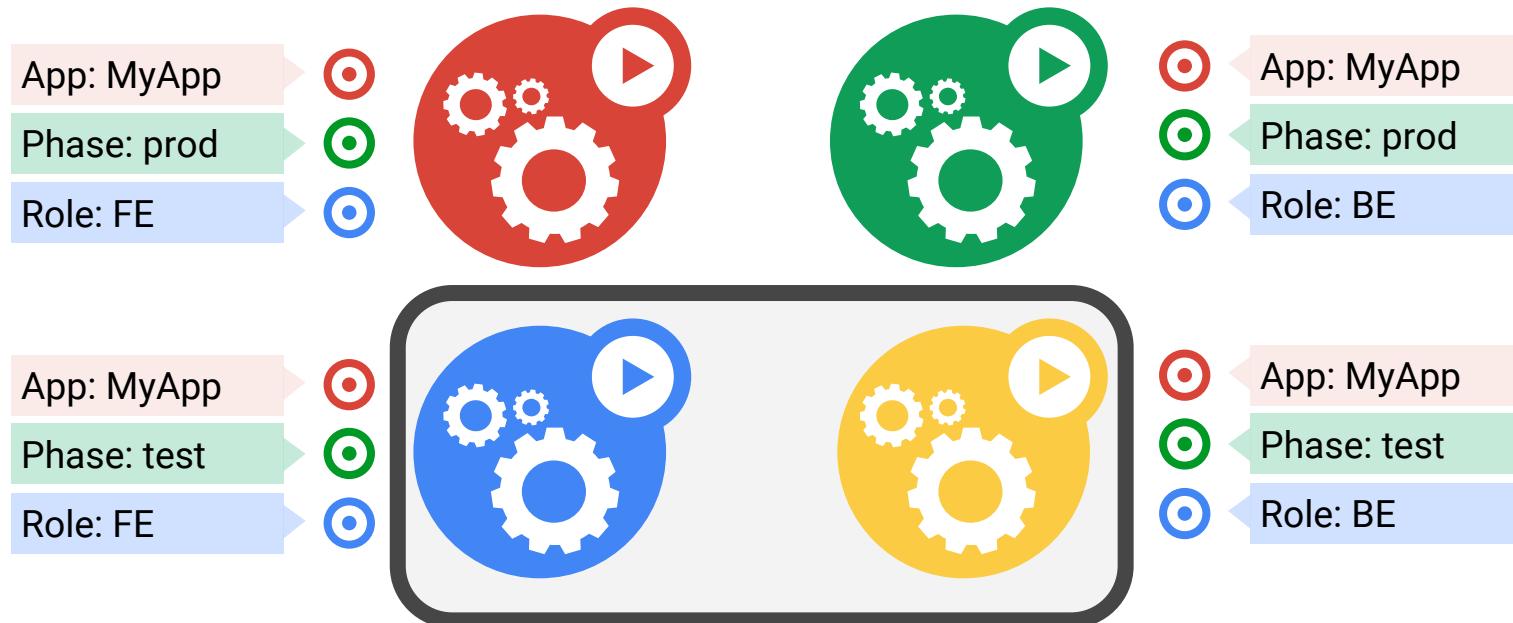
@meteatamel



App = MyApp, Phase = prod

Selectors

@meteatamel



App = MyApp, Phase = test

Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

ConfigMap

Selector

Secret



Resiliency & Redundancy



ReplicaSets*

@meteatamel

A simple control loop

Runs out-of-process wrt API server

One job: ensure N copies of a pod

- grouped by a selector
- too few? start some
- too many? kill some

Layered on top of the public Pod API

Replicated pods are **fungible**

- No implied order or identity

ReplicaSet

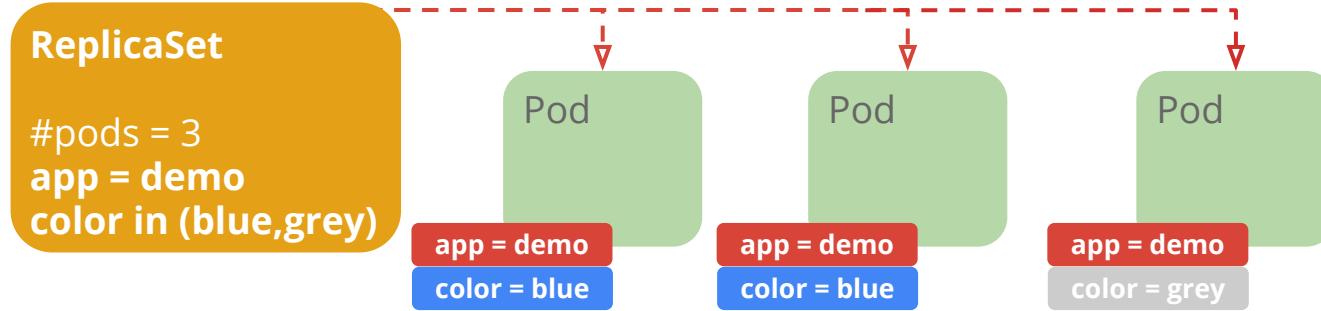
- **name** = "my-rc"
- **selector** = {"App": "MyApp"}
- **template** = { ... }
- **replicas** = 4



* The evolution of ReplicationControllers

ReplicaSets

@meteatamel



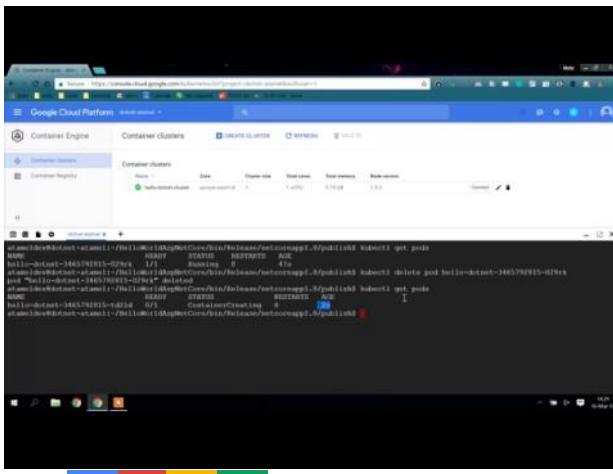
Behavior

- Keeps Pods running
- Gives direct control of Pod #s
- Grouped by Label Selector

Benefits

- Recreates Pods, maintains desired state
- Fine-grained control for scaling
- Standard grouping semantics

Demo: ReplicaSets



Kubernetes Health Checks



Health Check Philosophy



It's your responsibility to let Kubernetes know whether your app is healthy or not!

Liveness Probes

Liveness Probes make sure your application is running

```
livenessProbe:  
  # an http probe  
  httpGet:  
    path: /healthz  
    port: 8080  
  initialDelaySeconds: 15  # wait 15 seconds after pod is started to check for health  
  timeoutSeconds: 1        # wait 1 second for a response to health check
```



Readiness Probes

Readiness probes make sure your application is ready to serve traffic

```
readinessProbe:  
  # an http probe  
  httpGet:  
    path: /readiness  
    port: 8080  
  initialDelaySeconds: 20 # wait 20 seconds after pod is started to check for health  
  timeoutSeconds: 5      # wait 5 second for a response to health check
```

Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

ConfigMap

Selector

Secret



Services



A logical grouping of pods that perform the same function (the Service's endpoints)

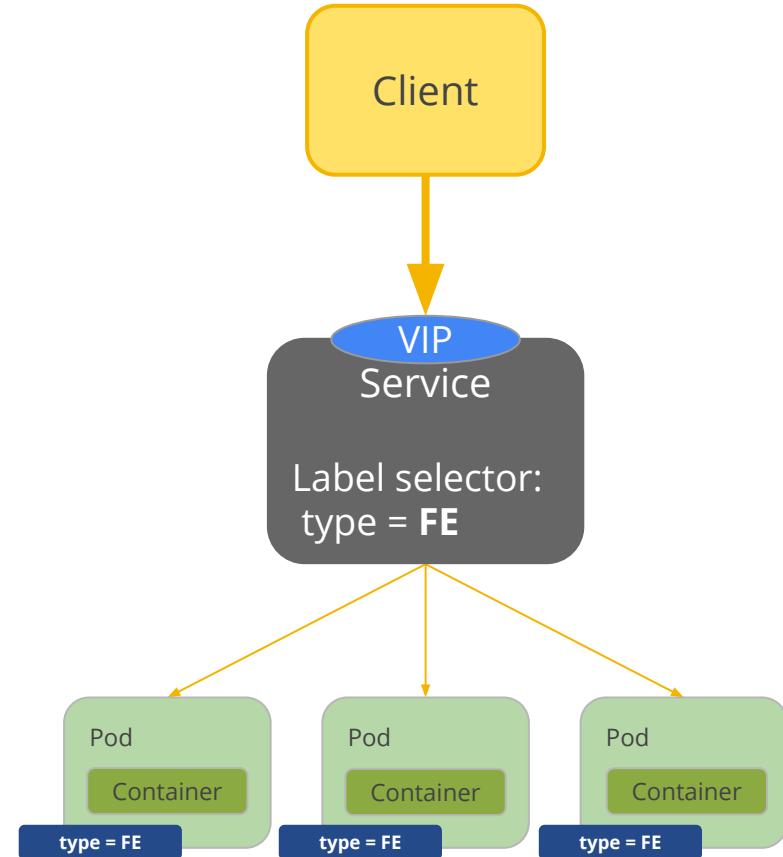
- grouped by label selector

Load balances incoming requests across constituent pods

Choice of pod is random but supports session affinity (ClientIP)

Gets a **stable** virtual IP and port

- also a DNS name



Demo: Services

The screenshot shows a Linux terminal window with several command-line outputs:

- Output from `kubectl get deployment`:

```
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
hello-world   1/1     1           1           10m
```
- Output from `kubectl get pods`:

```
NAME          STATUS    MESSAGE           AGE
hello-world-3663787151-471d4   Running   <none>            10s
```
- Output from `kubectl expose deployment hello-world --type=LoadBalancer`:

```
service "hello-world" exposed
```
- Output from `kubectl get services`:

```
NAME         CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
hello-world  10.3.25.118   <pending>    80/TCP        10s
```
- Output from `curl http://10.3.25.118:80`:

```
443/HTTP 200 1m
```
- Output from `kubectl get svc`:

```
NAME         CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
hello-world  10.3.25.118   <pending>    80/TCP        10s
```

Kubernetes Dashboard

A general purpose, web-based UI to view/manage Kubernetes clusters

The screenshot shows the Kubernetes Dashboard interface. On the left, there's a sidebar with navigation links: Admin (Namespaces, Nodes, Persistent Volumes), Namespaces (kube-system selected), Workloads (Deployments, Replica Sets, Replication Controllers, Daemon Sets, Stateful Sets, Jobs, Pods selected), and Services and discovery (Services). The main area has two charts: 'CPU usage' and 'Memory usage'. Below the charts is a table titled 'Pods' showing 1 - 10 of 18 results.

Name	Status	Restarts	Age	CPU (cores)	Memory (bytes)
dashboard SAME-i...	Running	0	14 minutes	0	5.4 Mi
fluentd-cloud-logg...	Running	0	2 months	0.01	120.8 Mi
fluentd-cloud-logg...	Running	0	2 months	0.009	94.3 Mi
fluentd-cloud-logg...	Running	0	2 months	0.014	174.3 Mi
heapster-v1.2.0-4...	Running	0	10 days	0.002	44.8 Mi

Demo: Kubernetes Dashboard

The screenshot shows the Kubernetes Dashboard interface. On the left is a sidebar with navigation links: Home, Workloads, Deployments, Services and Discovery, Storage, Config, and Help. The main area has two charts: 'CPU usage' and 'Memory usage'. Below the charts is a table titled 'Deployments' with one entry:

Name	Labels	Status	Age	Image
info-deployment	run=info-deployment	1/1	12 days	gcr.io/k8s-staging/hello-world

Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

ConfigMap

Selector

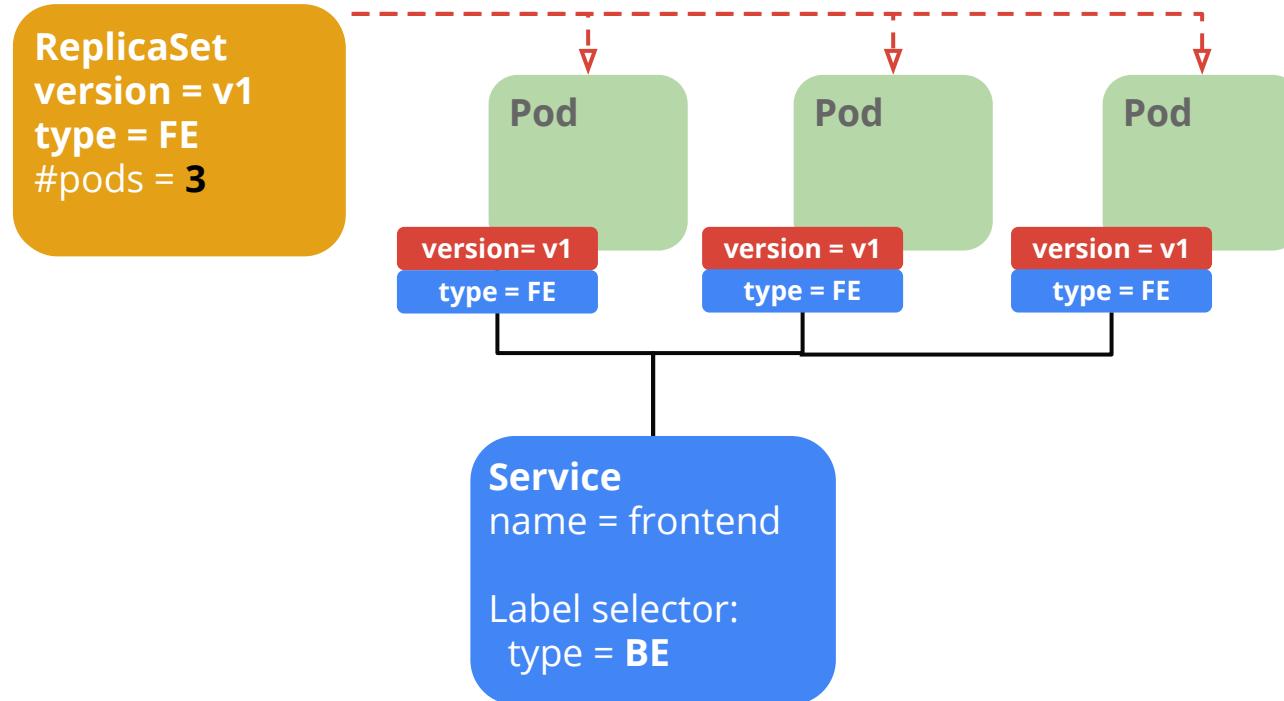
Secret



Scaling

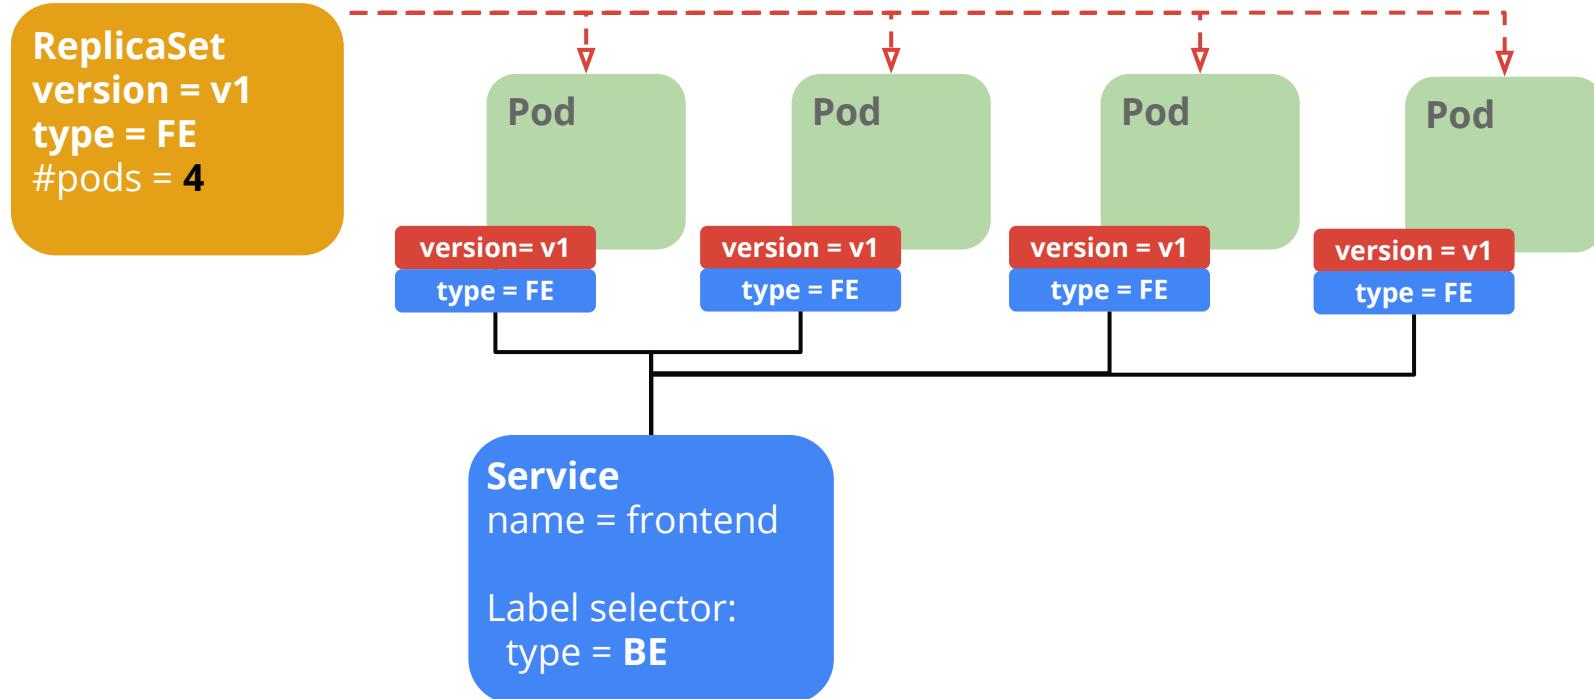
Scaling

@meteatamel

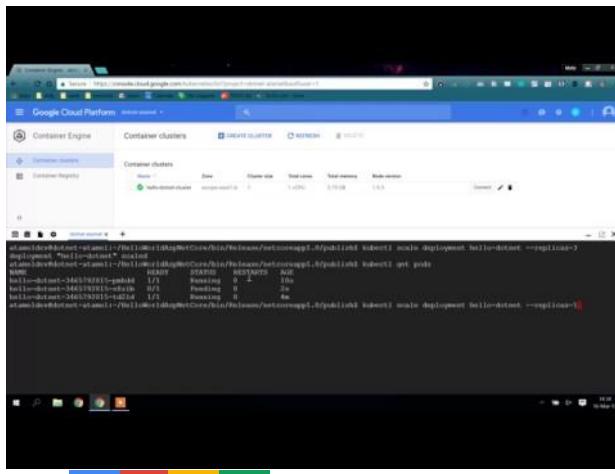


Scaling

@meteatamel



Demo: Scaling



The screenshot shows a laptop screen with two windows open. The top window is a browser displaying the URL <https://console.cloud.google.com/>. The bottom window is a terminal window with a dark background. The terminal output shows the deployment of a 'hello-distro' application using kubectl. The command entered was:

```
curl https://raw.githubusercontent.com/GoogleCloudPlatform/k8s-samples/main/deployments/hello-distro/hello-distro-deployment.yaml | kubectl apply -f -
```

The terminal then lists the deployment status:

NAME	READY	STATUS	RESTARTS	AGE
hello-distro-346378015-8d2b	1/1	Running	0	1m
hello-distro-346378015-8d2c	1/1	Running	0	4s

At the bottom of the terminal, there is a redacted line of text.

Rolling Update

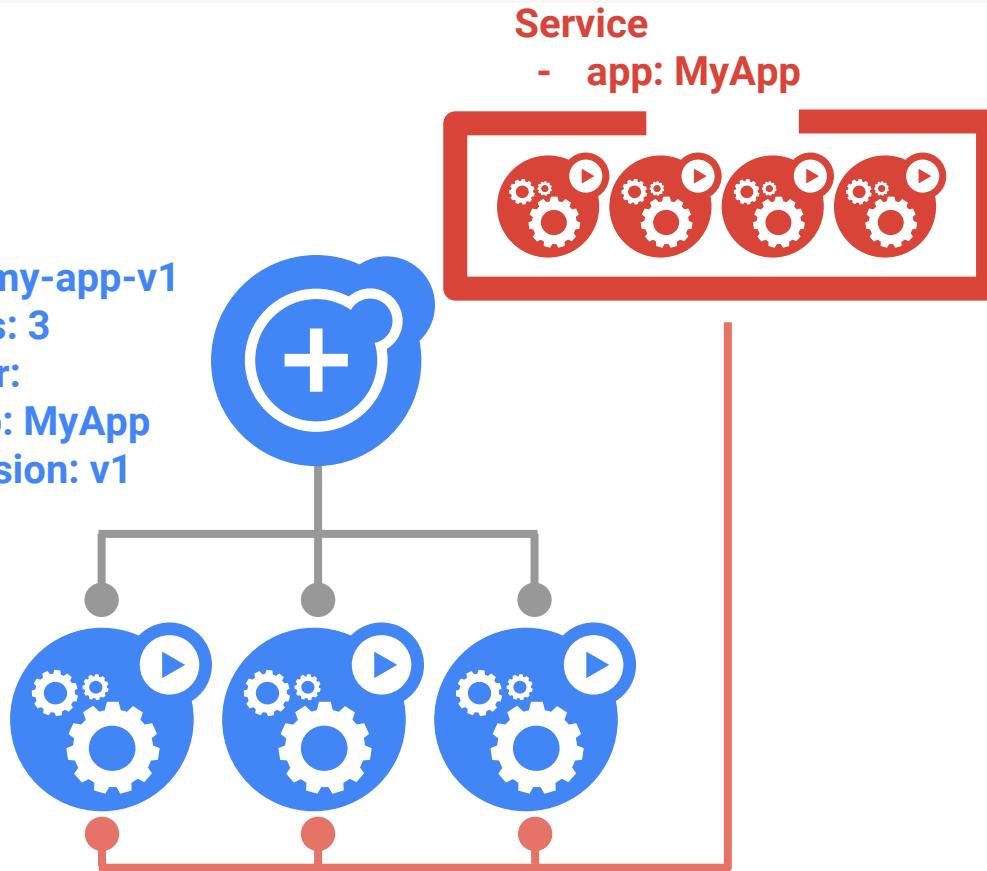


Rolling Update

@meteatamel

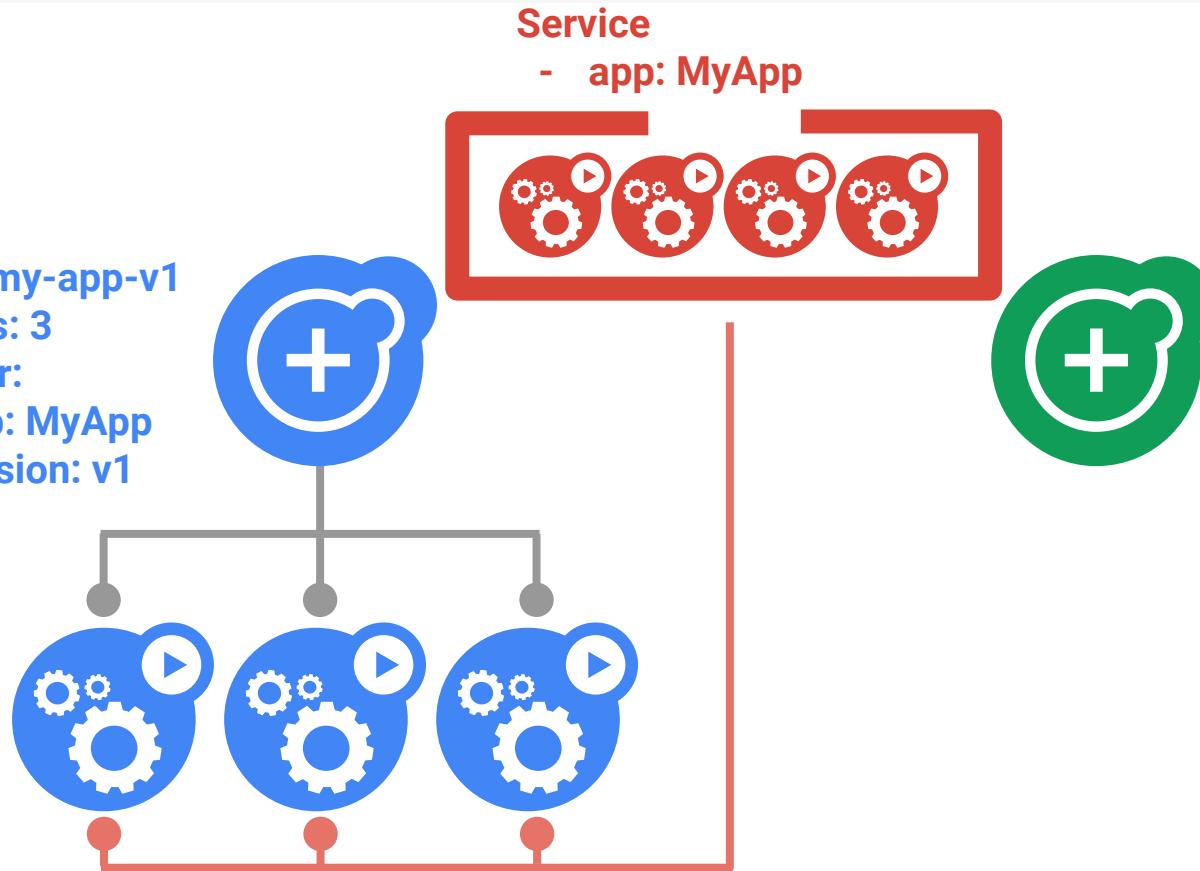
ReplicaSet

- name: my-app-v1
- replicas: 3
- selector:
 - app: MyApp
 - version: v1



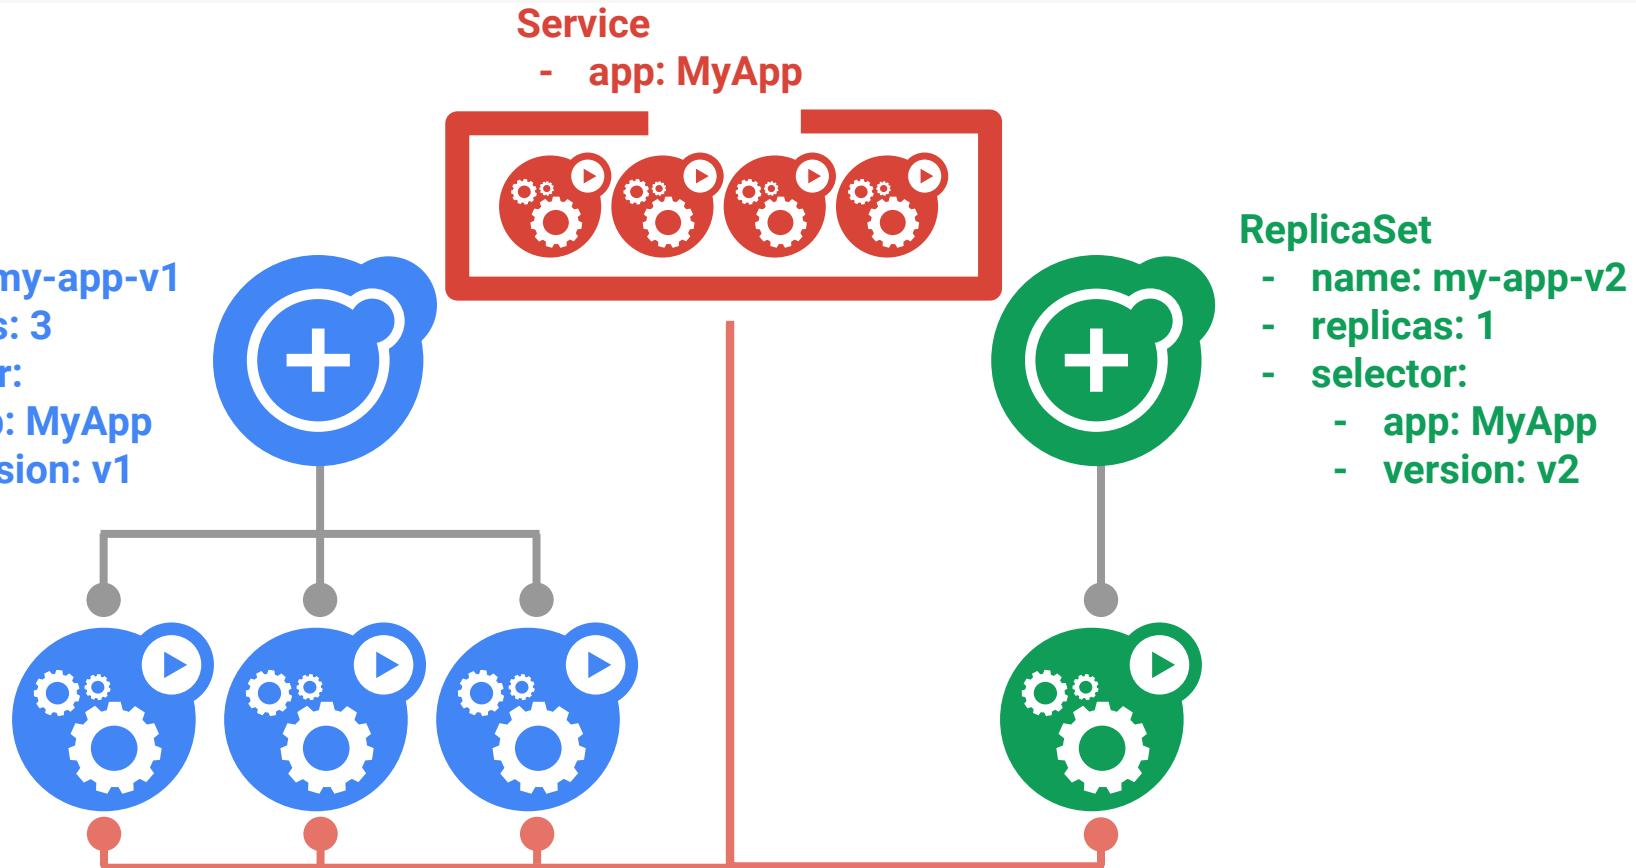
Rolling Update

@meteatamel



Rolling Update

@meteatamel

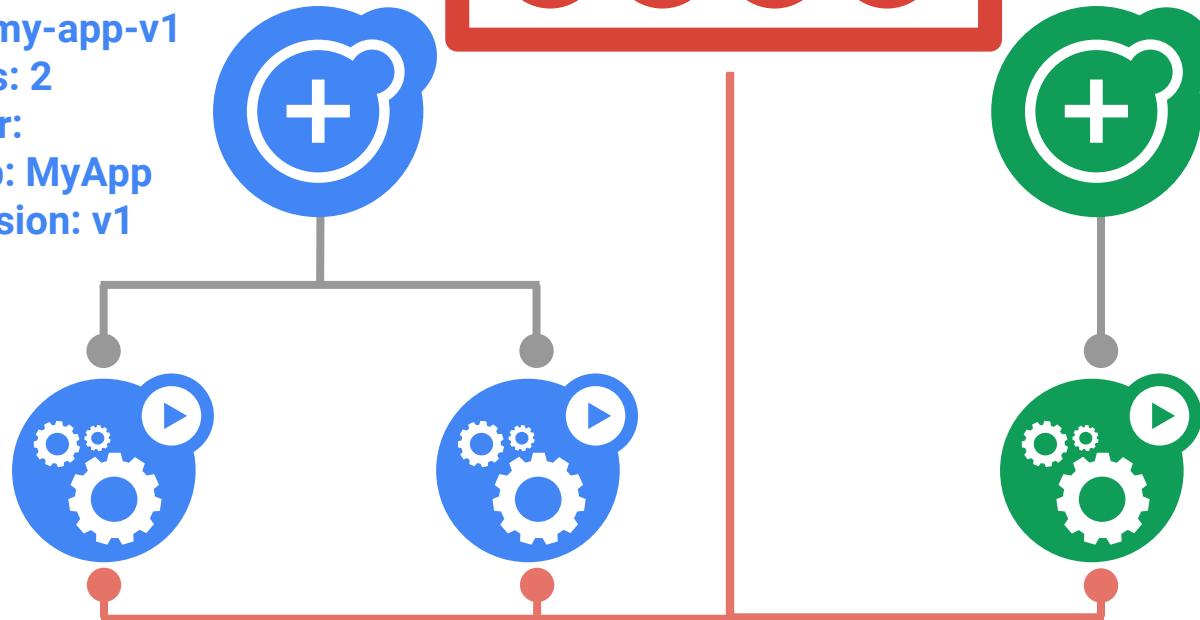


Rolling Update

@meteatamel

ReplicaSet

- name: my-app-v1
- replicas: 2
- selector:
 - app: MyApp
 - version: v1



ReplicaSet

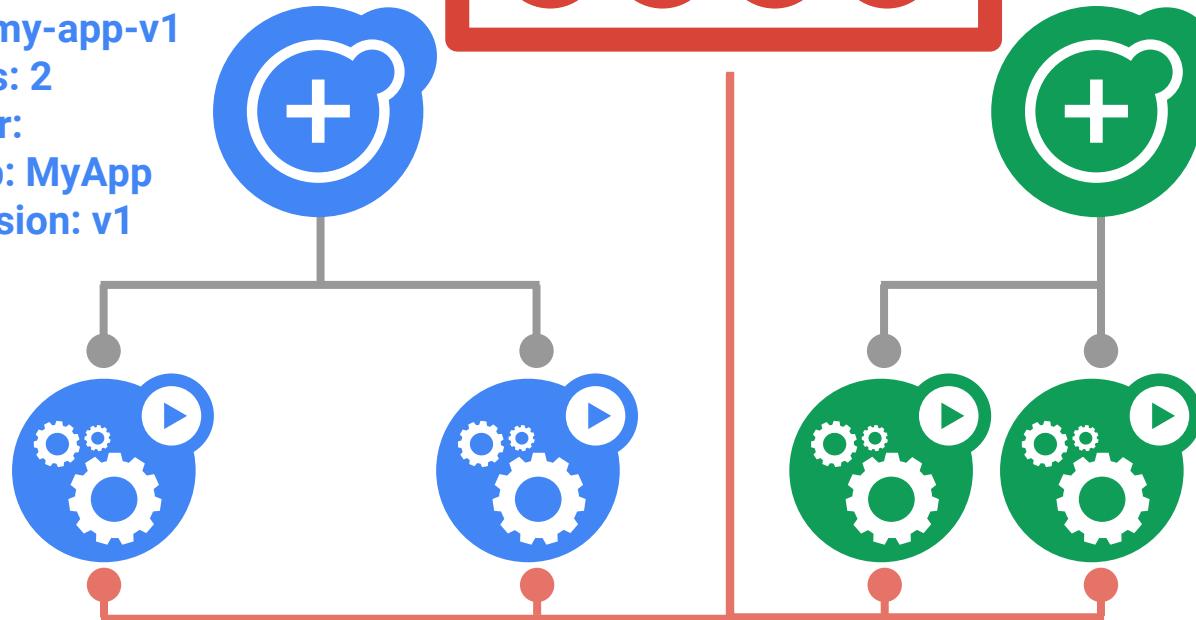
- name: my-app-v2
- replicas: 1
- selector:
 - app: MyApp
 - version: v2

Rolling Update

@meteatamel

ReplicaSet

- name: my-app-v1
- replicas: 2
- selector:
 - app: MyApp
 - version: v1



ReplicaSet

- name: my-app-v2
- replicas: 2
- selector:
 - app: MyApp
 - version: v2

Rolling Update

@meteatamel

ReplicaSet

- name: my-app-v1
- replicas: 1
- selector:
 - app: MyApp
 - version: v1



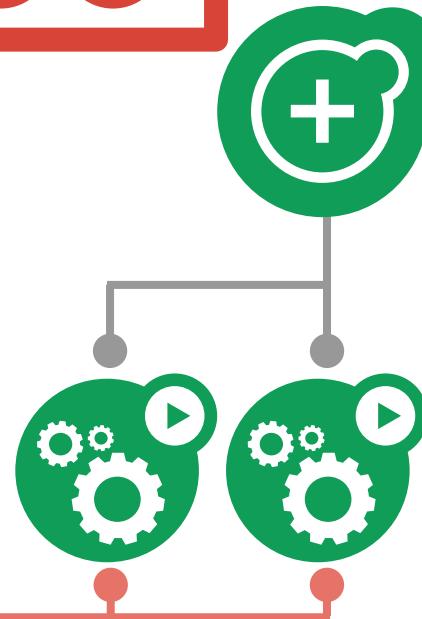
Service

- app: MyApp



ReplicaSet

- name: my-app-v2
- replicas: 2
- selector:
 - app: MyApp
 - version: v2



Rolling Update

@meteatamel

ReplicaSet

- name: my-app-v1
- replicas: 1
- selector:
 - app: MyApp
 - version: v1



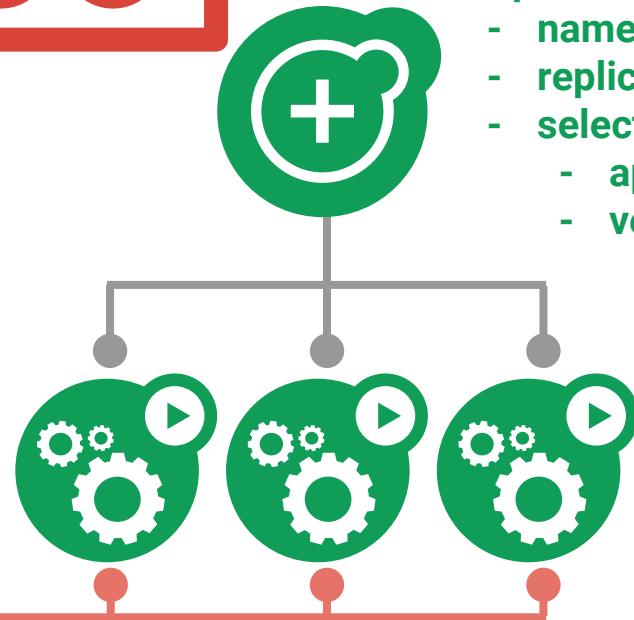
Service

- app: MyApp



ReplicaSet

- name: my-app-v2
- replicas: 3
- selector:
 - app: MyApp
 - version: v2



Rolling Update

@meteatamel

ReplicaSet

- name: my-app-v1
- replicas: 0
- selector:
 - app: MyApp
 - version: v1

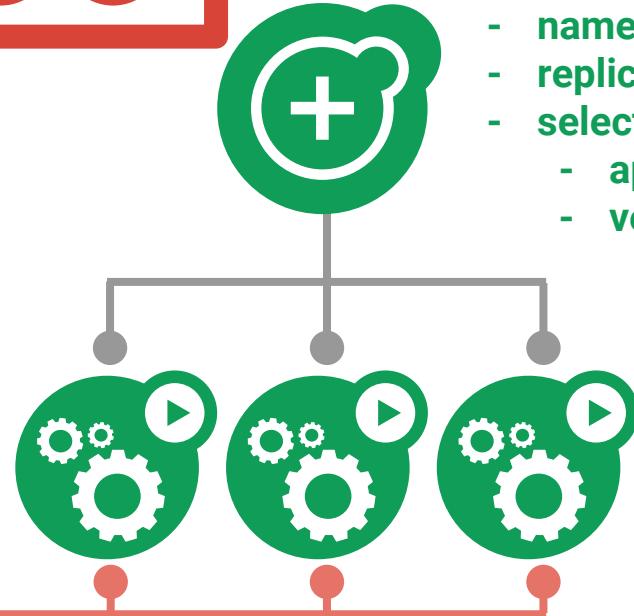


Service
- app: MyApp



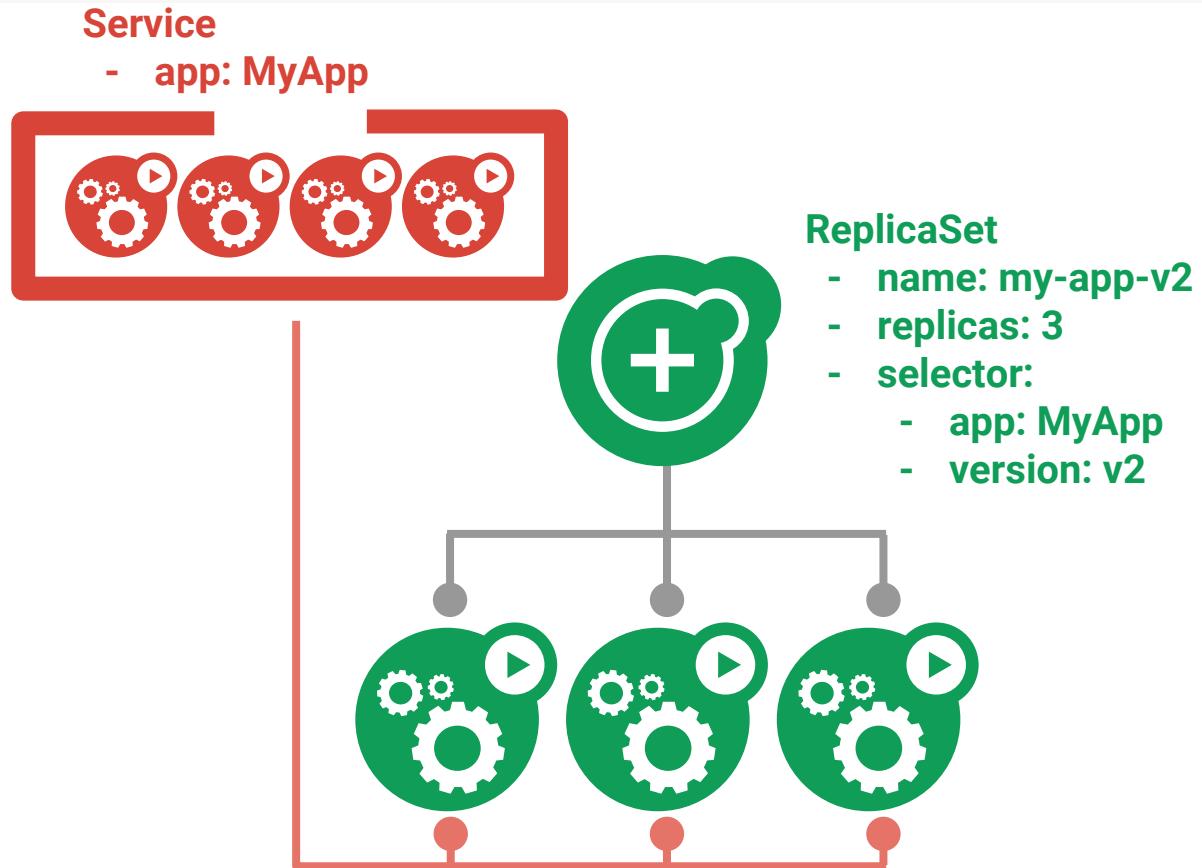
ReplicaSet

- name: my-app-v2
- replicas: 3
- selector:
 - app: MyApp
 - version: v2

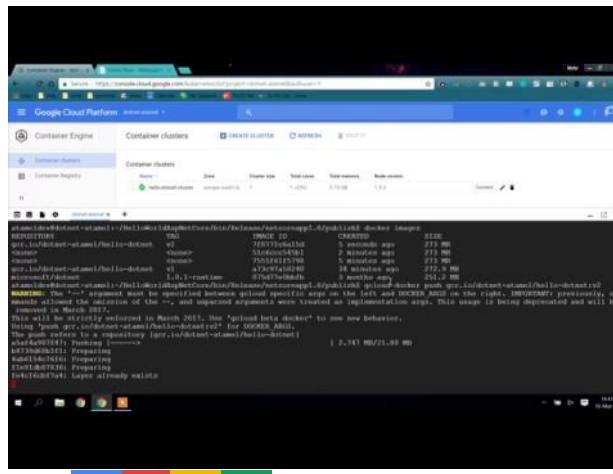


Rolling Update

@meteatamel

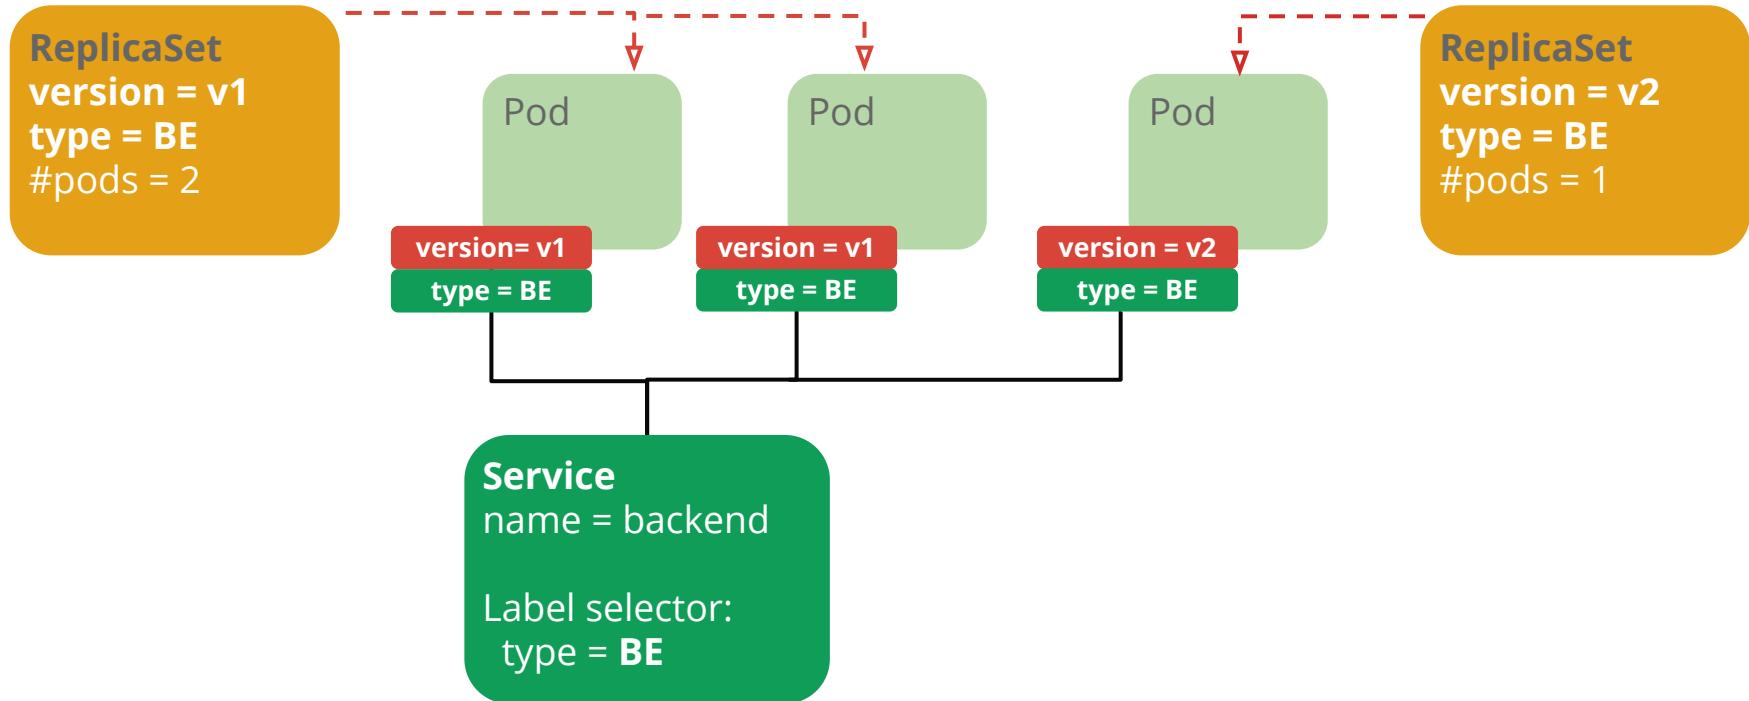


Demo: Rolling Update



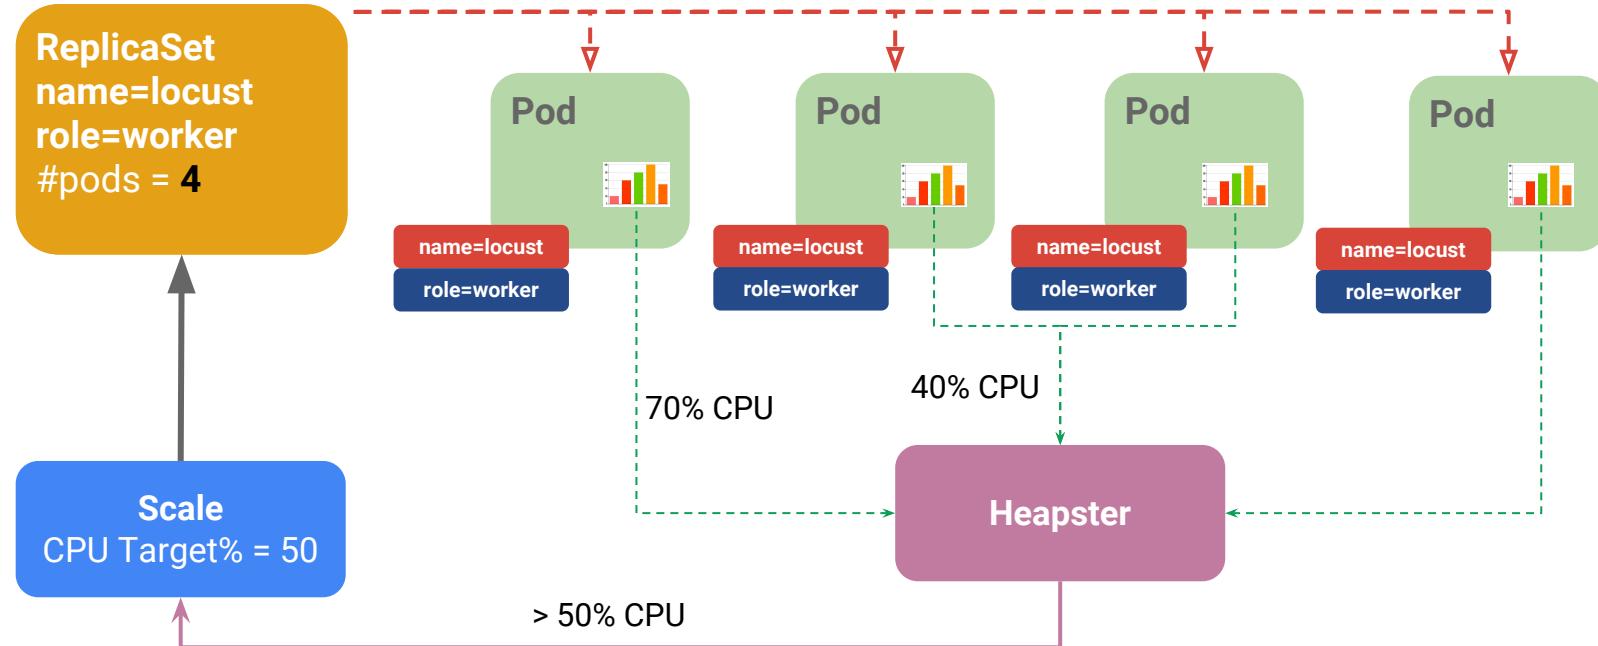
Canary Deployments

@meteatamel



Autoscaling

@meteatamel



DaemonSets



Problem: how to run a Pod on every node?

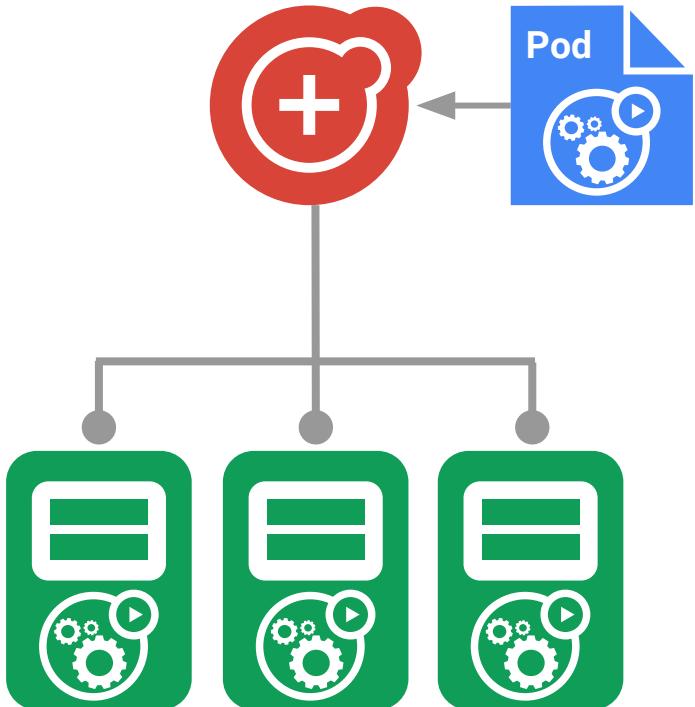
- or a subset of nodes

Similar to ReplicaSet

- principle: do one thing, don't overload

“Which nodes?” is a selector

Use familiar tools and patterns



Jobs

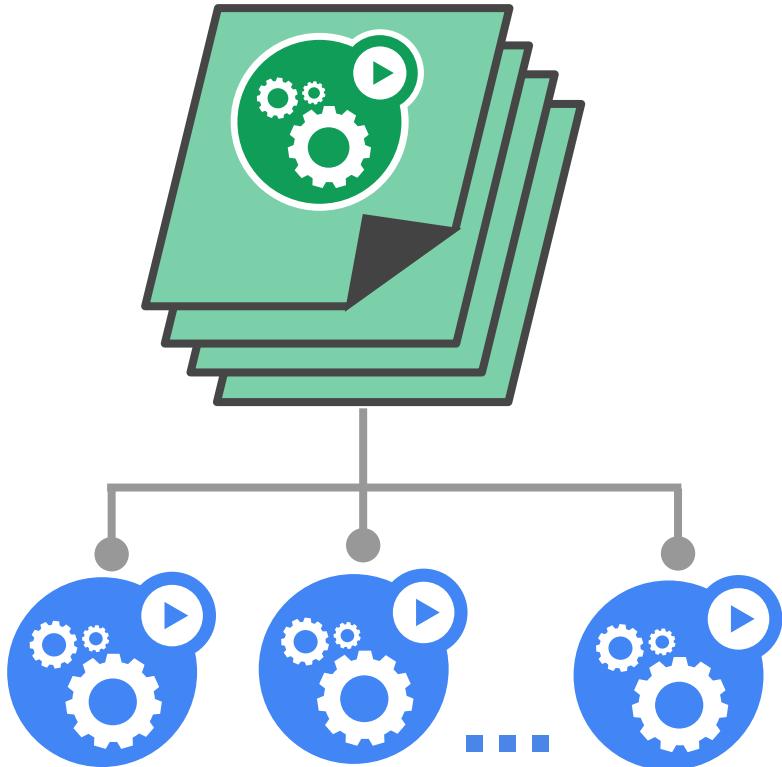


Run-to-completion, as opposed to run-forever

- Express parallelism vs. required completions
- Workflow: restart on failure
- Build/test: don't restart on failure

Aggregates success/failure counts

Built for batch and big-data work



StatefulSets

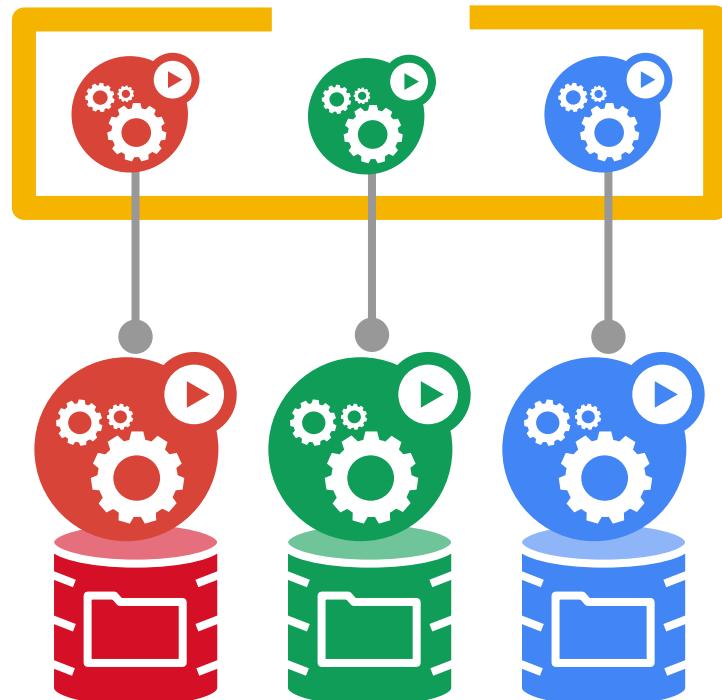


Goal: enable clustered software on Kubernetes

- mysql, redis, zookeeper, ...

Clustered apps need “identity” and sequencing guarantees

- stable hostname, available in DNS
- an ordinal index
- stable storage: linked to the ordinal & hostname
- discovery of peers for quorum
- startup/teardown ordering



ConfigMaps



ConfigMaps

@meteatamel

Goal: manage app configuration

- ...without making overly-brITTLE container images

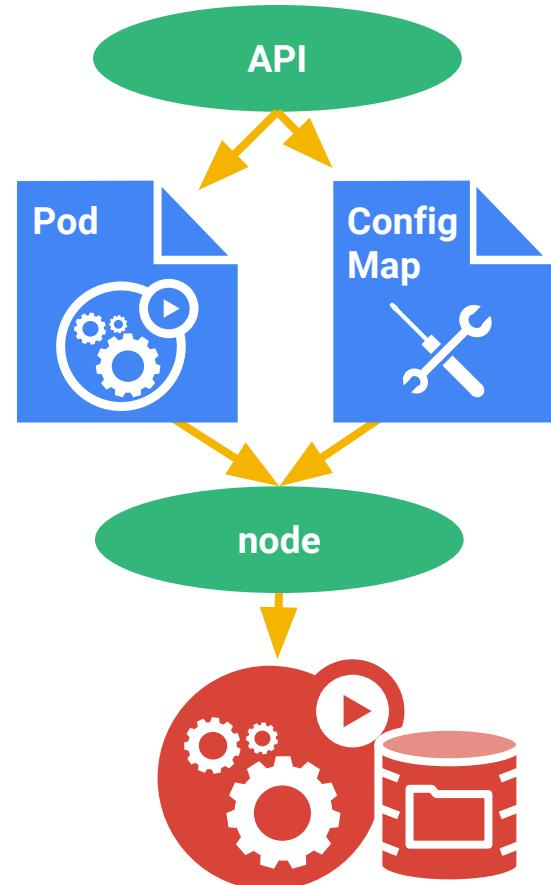
[12-factor](#) says config comes from the environment

- Kubernetes is the environment

Manage config via the Kubernetes API

Inject config as a virtual volume into your Pods

- late-binding, live-updated (atomic)
- also available as env vars



Secrets

Goal: grant a pod access to a secured *something*

- don't put secrets in the container image!

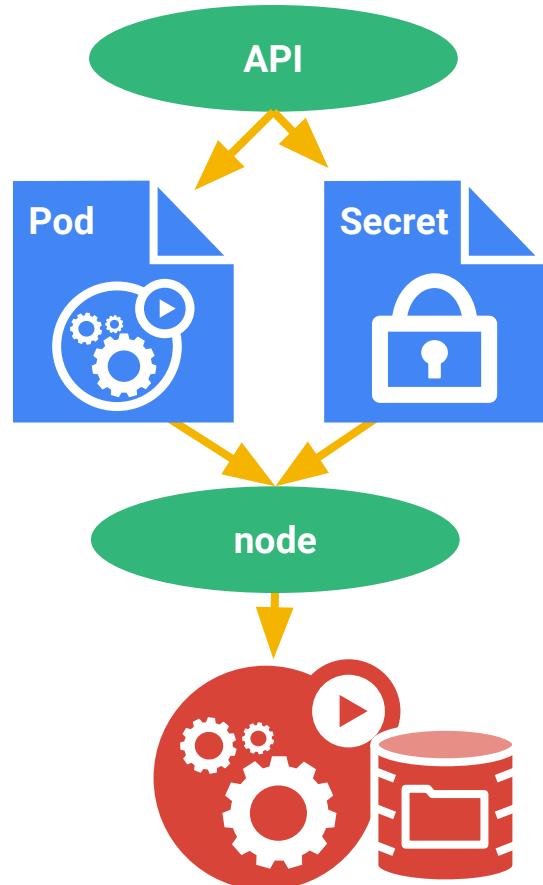
[12-factor](#) says config comes from the environment

- Kubernetes is the environment

Manage secrets via the Kubernetes API

Inject secrets as virtual volumes into your Pods

- late-binding, tmpfs - never touches disk
- also available as env vars



Kubernetes Terminology

Deployment

ReplicaSet

DaemonSet

Pod

Liveness Probe

Job

Volume

Readiness Probe

StatefulSet

Label

Service

ConfigMap

Selector

Secret



There is more!



Thank You

Send talk feedback
bit.ly/atamel

kubernetes.io
cloud.google.com/container-engine



Mete Atamel
@meteataamel
atamel@google.com
meteataamel.wordpress.com